Customer touchpoints and mobility service value

A proposed service model connects consented intent to reachable supply, additional activity and feasible travel choices.

01 / PRODUCT OUTCOME & SCOPE

PRODUCT HYPOTHESIS

Customer intent can improve service capacity and net contribution through three testable mechanisms.

Resident / visitor

Share plans, free time and preferences with consent. Retain control of timing, price and accessibility needs.

Mobility operator

Position reachable vehicles, protect committed trips and own exceptions within finite operating capacity.

Service provider

Return activity slots, prices and service capacity. Confirm reservations and observed use.

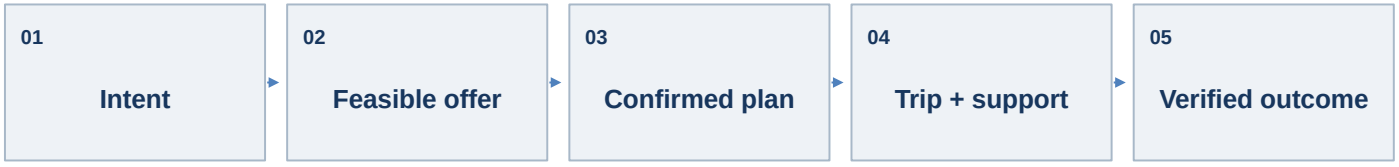
City partner

Set ODD, access and coverage conditions. Review outcomes for underserved places and users.

MaaS product layer

C01 identity / API → C02 booking & consent → C03 feasibility / reservation → C04 dispatch & status → C06 support / cases → C07 usage / settlement | C05 energy capability constrains feasible supply

END-TO-END SERVICE LOOP



THREE VALUE MECHANISMS / INPUT, ACTION AND TEST

01 Advance positioning

Received trip plans inform earlier positioning. Test pickup time, empty kilometres and unmet requests under the same demand and fleet.

02 Additional demand

Free time and preferences inform feasible activity offers. Separate genuinely additional use from trips or spending merely transferred.

03 Demand smoothing

Activity, travel time and total price form one offer. An accepted time change moves the same request to a less congested interval.

OPERATING BOUNDARIES

All effects depend on consent, accepted terms and finite vehicle, hub and provider capacity. Extra demand can worsen congestion. Positioning can add empty travel. Discounts can erase margin. Compare each mechanism and their combination, including refused offers, cancellations and recovery costs.

SERVICE CONTRACT C02 owns intent / booking revisions. C03 validates resource holds. C04 records execution. C06 owns exceptions. C07 reconciles actual use and costs. C05 constrains energy feasibility. Operating conditions: P25–P28.

End-to-end journey across user, operation and platform

The same intent remains traceable across recommendations, time changes, physical positioning and observed service outcomes.

02 / SERVICE BLUEPRINT



STAGE EVIDENCE **Plan: consent + receipt time** **Request: eligible offer + capacity** **Confirm: accepted terms + committed holds** **Travel: observed pickup + owned exception** **Complete: trip / activity actuals + cost**

EXCEPTION LOOP

A revised plan releases obsolete holds only after reconciliation. Cancellation or capacity loss updates positioning, ETA and the recovery owner.

TRACEABILITY CONTRACT

intentId + offerId + experimentArm + bookingRev + resource bundle + missionId

Record offeredAt, receivedAt, acceptedAt, original / revised time, price, consentRef and source revision. Link every time change to its original intent. C02 owns booking, C03 feasibility, C04 execution and C06 case lifecycle.

ACCEPTANCE INVARIANTS

Positioning uses only information received by the decision time. A suggestion never creates a booking without acceptance. Commit vehicle and activity holds before confirmation. Time changes preserve one request identity. Duplicate retries return the prior result. Stale capacity blocks a firm offer. Unknown command results require reconciliation. Completion needs observed actuals. Every open exception retains an owner and deadline.

MVP boundary and product prioritization

A complete service slice provides the baseline. Three separately controlled mechanisms then test customer and operating value.

03 / MVP & PRIORITIZATION

Priority rule

Rank by measured service benefit, net contribution, feasibility and learning value.

Common fleet / hubs / demand conditions support comparison. Capacity and service guardrails remain binding.

PILOT MVP

One selected service slice

01	Capture intent, needs, consent and idempotent request
02	Return feasible options, binding reasons and expiry
03	Commit booking + resource bundle with revision checks
04	Track state, freshness, support case and recovery owner
05	Reconcile actual trip, service use and passenger outcome

Exit only when normal + recovery paths pass acceptance; operator can resolve exceptions; event and cost records reconcile.

ITERATE

Three controlled mechanisms

01	Use received plans for physically feasible positioning
02	Offer activities within spare time and provider capacity
03	Offer accepted time changes with transparent total price
04	Keep independent switches and log combined interactions
05	Separate new demand, shifted intent and displaced use

Advance only when gains survive service guardrails, realistic acceptance, full costs and adverse demand / supply conditions.

EXTEND PLATFORM

Repeatable operating model

01	Versioned adapters and purpose-limited consent
02	Energy, hub throughput and staffing in feasibility
03	Guarded support and owned passenger recovery
04	Calibrate by zone, time and user segment
05	Reusable configuration, evidence and rollback

Expand after second-site tests preserve service quality and net contribution with local demand, capacity and operating costs.

MVP CONTROL

MVP entry criteria: selected geography / ODD / channel / segment; named service and incident owners; provider, fleet and site capacity confirmed; legal, privacy, accessibility and partner constraints reviewed; C01–C07 data and authority boundaries agreed. Pilot exit: normal + cancellation + disruption / recovery traces pass; queue and support coverage are staffed; KPI baseline, missing-data behavior and cost ledger reconcile; rollback and customer communication are rehearsed. Optional C05 optimization, remote assistance and I09 roadside sensing stay out until separately gated.
G2 constraint register: ODD = approved operating conditions (road, speed, weather, time, etc.). DDT = real-time operational / tactical functions needed to operate in traffic (including OEDR); record which actor performs them. Define DDT fallback separately by actor, trigger, response and intervention boundary. Also record operating limits (hours / staffing / recovery), infrastructure (node power / energy / footprint) and business limits (fare / demand / transit overlap). For each, record variable, unit / value or TBD, source URL, owner, affected KPI and PL line. Set thresholds only after evidence and ownership are confirmed. Source: SAE J3016_202609, doi:10.4271/J3016_202609.

Interaction concept for user app and operator workspace

A simple request–offer–service flow exposes meaningful choices. Detail appears when the user compares or changes a plan.

04 / UX FLOW & PROTOTYPE

RESIDENT / VISITOR APP • CONCEPTUAL

01 REQUEST	02 OFFER	03 SERVICE
Plan a journey	**Choose a feasible plan**	**Track and get help**
Origin / destination + pickup point	Activity + mobility availability	Booking / mission state + age
Time window + flexibility	Original / alternative time + expiry	ETA + last update
Party size + access / luggage	Access fit + pickup / arrival	Pickup / onboard / arrival proof
Optional plan / preference consent	Total price + change / refund terms	Help, cancel or recovery path
Validate → submit intent	**Accept plan / keep original**	**Confirm outcome / recover**

OPERATOR WORKSPACE • ACTION CONTEXT

CAPACITY & QUEUE

Demand by zone / time
Reachable vehicles + ready time
Activity / hub slots and queues
Committed trips + spare capacity
Cases by severity / dueAt
Owner + recovery availability

WHY THIS OFFER

Received plan + consent scope
Current capacity / evidence age
Original vs revised trip timing
Positioning distance + cost
Discount + expected net contribution
UNKNOWN remains explicit

PERMITTED ACTION

Position within current limits
Hold / release finite resources
Propose terms, await acceptance
Retain qualified case owner
Reconcile timeout / lost ACK
Audit every accepted change

USABILITY TASKS TO TEST

Test whether users understand optional data access, the original choice, changed timing and total cost. Measure task success, critical errors, opt-out and acceptance by accessibility need. Include rejected suggestions, stale capacity, cancellation after positioning and lost acknowledgement. Operators must explain infeasibility, protect confirmed trips and recover without duplicate bookings.

Prototype evidence: decision time, price / timing comprehension, recovery completion and reasons for refusal. No forced consent or automatic schedule change.

STATE MODEL

DRAFT → QUOTED (expiresAt) → PENDING (hold) → CONFIRMED (commit) → IN_SERVICE → COMPLETED (actuals + passenger outcome) / CANCELLED (reason + release). Timeout / lost acknowledgement → UNKNOWN → query and reconcile; never duplicate a side effect. Revision conflict → refresh terms; stale evidence → block offer or action with reason.

Three mechanisms, service guardrails and net value

Measure causal effects against a matched baseline. Keep demand identity, service denominators and cost boundaries explicit.

05 / SUCCESS MEASURES

NORTH STAR

Verified completed journeys

count only when mission evidence + service actual + user outcome reconcile

01 POSITIONING

Pickup wait

Actual pickup − requested pickup time. Report median / p95.

C02 request + C04 arrival

Empty travel

Empty vehicle-km / completed trip. Include advance positioning.

C04 movement trace

Dispatch success

Assigned missions / eligible requests. Retain capacity rejects.

C02 requests + C04 assignments

Required fleet

Minimum tested fleet meeting all service gates, with uncertainty.

Fleet sweep + P25–P28

02 NEW DEMAND

Offer conversion

Accepted activity offers / eligible offers shown.

C02 offer and acceptance

Additional trips

Treatment − matched baseline completed trips. Identify substitution.

Intent linkage + comparison arm

Fleet utilization

Productive / available hours. Also report scheduled hours.

C04 + fleet schedule

Activity use

Verified additional provider visits. Keep gross and net use separate.

Provider actuals + comparison

03 SMOOTHING

Shift acceptance

Accepted time changes / eligible shift offers shown.

C02 offer + booking revision

Peak demand

Requests per fixed interval before / after accepted time shifts.

Original and revised request time

Unserved demand

Unserved eligible requests / all eligible requests by zone and time.

C02 + final request disposition

Promise kept

On-time pickups / pickups due. Include missed pickups.

Accepted window + C04 arrival

NET ECONOMICS

Net revenue

Recognized revenue after discounts / refunds, counted once.

C07 contract + ledger

Incremental cost

Extra travel, energy, service and app operating cost.

C04 / C05 + cost ledger

Incremental contribution

Δ recognized revenue − Δ attributable cost. Report total and per completed trip. Include app cost once and document allocation.

Matched arms + C07 ledger

Interaction effect

Combined gain − sum of separately tested gains. Report overlap and capacity effects.

Paired runs / controlled pilot

COMMON GUARDRAILS

Completion, pickup promise, recovery and coverage

Keep unmet demand, support load, energy / hub queues and user opt-out visible. Check outcomes by zone, peak interval and accessibility need.

MEASUREMENT CONTRACT

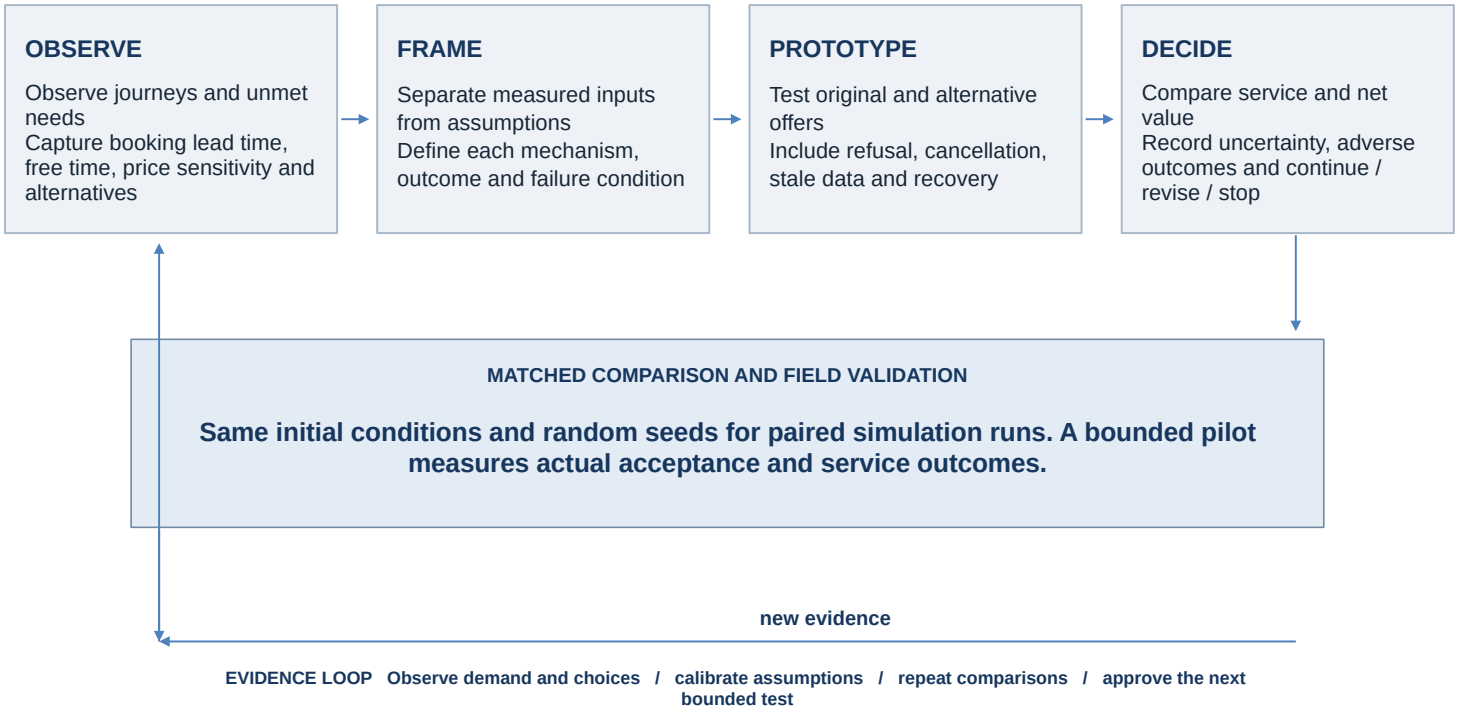
No effect size or target is asserted. Set thresholds from an agreed baseline, then validate in simulation and a bounded field pilot.

ATTRIBUTION: classify additional demand, changed-time demand and substituted use separately. A shifted request keeps its identity and does not increase the request total. Compare the same initial fleet, hubs, demand trace, outages and cost basis. Report repeated-run uncertainty and missing events. Combined effects need explicit interaction tests.
SERVICE: completion = completed trips / eligible requests. Recovery = recovered / eligible disrupted journeys. Available vehicle-hours and scheduled hours remain separate. Include rejected and expired requests in demand reporting. Capacity, replenishment and degraded operation: P25–P28.
FINANCE: use one revenue boundary and one cost ledger. Deduct discounts once. Distinguish operator revenue from provider revenue and user spending. Report total contribution alongside per-trip contribution. Keep CAPEX, operating cash flow, payback and IRR separate. Each input retains unit, source, owner and calculation link.

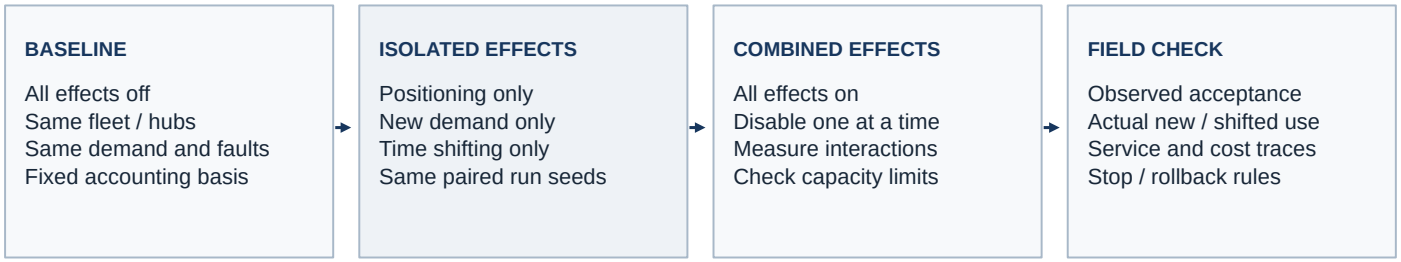
Controlled experiments and evidence workflow

Simulation tests the mechanism under stated assumptions. Field observation tests whether users respond and operations reproduce the effect.

06 / DISCOVERY & VALIDATION



EXPERIMENT DESIGN / ATTRIBUTION AND INTERACTION



STRESS TESTS

Sweep booking lead time, forecast error, acceptance, additional demand and discount level. Test peaks, spatial imbalance, hub / vehicle loss, replenishment queues and cancellation after positioning. Reject future-information leakage. Preserve shifted request identity and record unmet demand. A gain that disappears with realistic costs or breaches service limits fails the gate. AI may assist analysis, but owners verify evidence and retain operational authority.

Proposed tests only. Simulation outputs do not establish field impact or guarantee service-level attainment.

REPRODUCIBLE EVIDENCE RECORD

Record scenario and engine version, input units / source, demand trace, received-information timestamps, initial fleet / hub state, random seed and effect switches. Retain offer / intent linkage, acceptance, original and revised time, trip / activity actuals, rejected requests and cost entries. Report paired differences and uncertainty across repeated runs.
Field design: choose assignment unit and observation window in advance. Account for spillover across shared vehicles and zones. Measure non-participation, displacement and rebound demand. Name the decision owner, acceptance thresholds, adverse-outcome limits and rollback trigger.

Roadmap from discovery to repeatable deployment

Each gate adds evidence for customer response, operating feasibility and net value before the service expands.

07 / RELEASE & SCALE GATES

G1 / DISCOVER

Select users, zones and time windows. Observe lead time, unmet demand, activity alternatives and willingness to shift. Establish traceable service and cost baselines. Record sampling gaps and consent boundaries.

EVIDENCE Journey / demand map, request denominators, current alternatives and observed constraints

GO / HOLD / REVISE

G2 / DEFINE

Define the three mechanisms and separate controls. Prototype original and alternative offers. Set acceptance, cost boundaries and service guardrails. Agree the comparison design, sources and owners before evaluating results.

EVIDENCE Causal model, offer tests, KPI dictionary, baseline / treatment definition and stop rules

GO / HOLD / REVISE

G3 / PILOT READY

Pass normal and failure paths for C01–C07. Verify signal receipt timing, finite holds, no duplicate shifted demand and auditable costs. Staff support and recovery. Rehearse outages, cancellation, rollback and customer communication.

EVIDENCE Acceptance traces, capacity / energy checks, accounting reconciliation and operating readiness

GO / HOLD / REVISE

G4 / LEARN

Compare individual and combined effects with matched baselines. Validate actual acceptance and additional use. Check peak / zone service quality, displacement, interaction and net contribution. Re-test uncertain inputs before expanding.

EVIDENCE Effect range, adverse outcomes, full costs, counter-evidence and continue / revise / stop decision

GO / HOLD / REVISE

G5 / REPEAT

Repeat at a second site with local demand, provider capacity and costs. Confirm gains survive weak acceptance and supply loss. Package configuration, training, monitoring and recovery. Approve expansion within a verified operating envelope.

EVIDENCE Second-site results, sensitivity, deployment / support cost, partner acceptance and rollout owner

GO / HOLD / REVISE

HOLD / REPLAN RULES

Schedule alone never passes a gate. HOLD for unresolved safety / privacy / legal risk, unknown execution, missing evidence, unstaffed support, partner unreadiness, unreliable data or unreconciled cost.

REVISE with named owner, required evidence, deadline and temporary scope restriction. Return to an earlier gate when new evidence changes the user problem or service boundary. Record rollback trigger and resume authority.

GATE TRANSITION RULE • Gates are evidence checkpoints, not fixed calendar phases. A failed condition returns to the earliest gate whose assumption is invalidated; scope stays held until the named owner closes the evidence gap and the decision is recorded.

PRODUCT DECISION OWNERSHIP

PM: user outcome, scope, priority and KPI decision | Design: interaction and usability | Engineering: architecture and delivery | Operations / partners / city: service readiness and constraints

TECHNICAL EVIDENCE IN THE FOLLOWING SYSTEM DESIGN

Reservation and service contracts: slides 8–11 • Remote operation and roadside interfaces: slides 13–23 • Fault injection and release dependencies: slide 24

Operating capacity, replenishment, recovery and verification conditions follow on slides 25–28. All tests remain proposed.

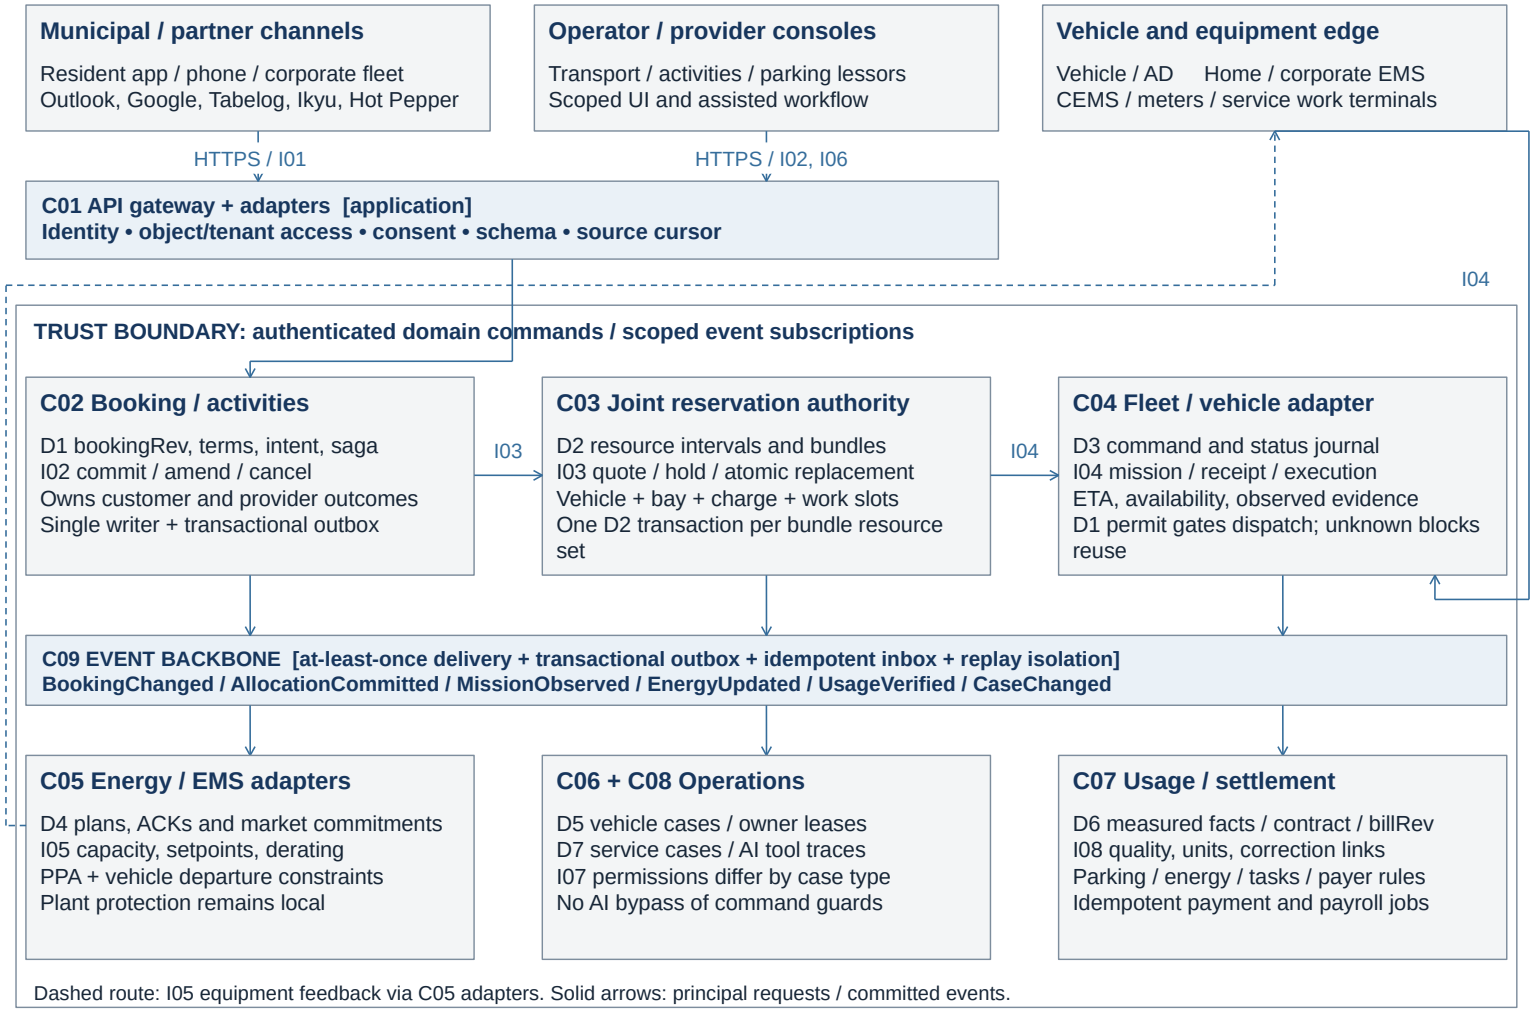
GO / HOLD RECORD

Decision + date • user segment / site • baseline and observed KPI • acceptance trace • unresolved safety / legal / operational risk • scope restriction • accountable owner • next review trigger. GO requires named service owner, trained operation, support coverage, data quality and rollback path. Finance record: CAPEX and OPEX by node and project; node and project payback; IRR; sensitivity range; Variable / Value / Source URL / formula-use link for every assumption.

Container architecture and deployment boundaries

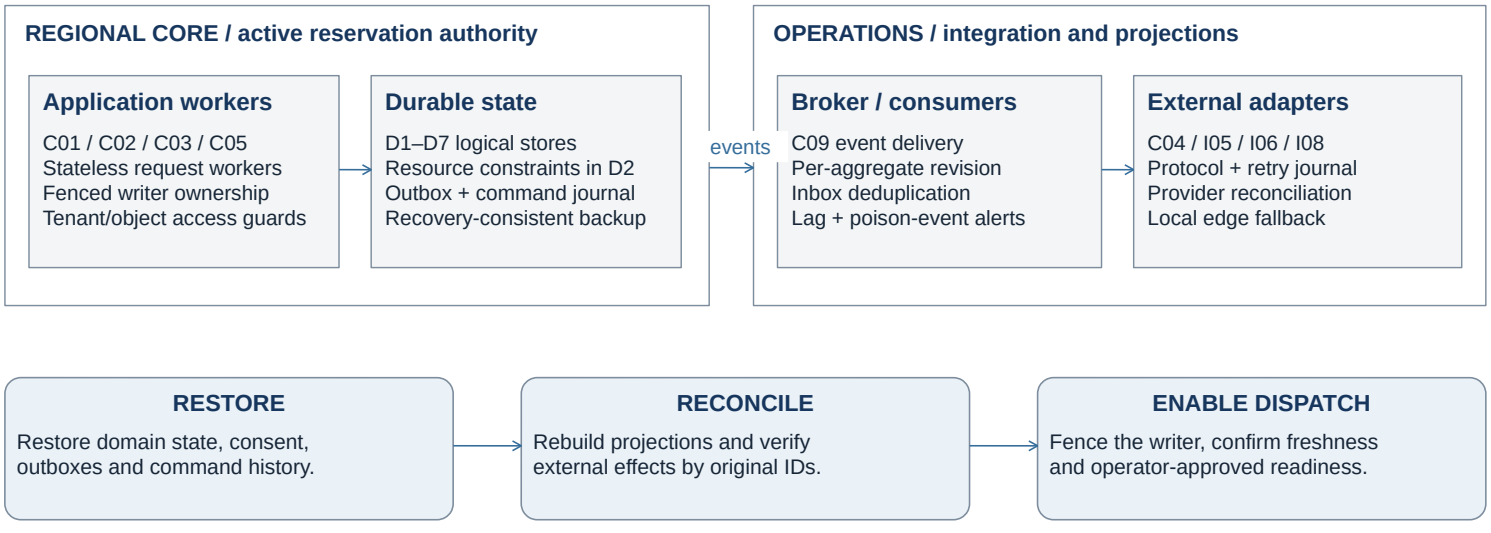
Proposed common PaaS for integrated reservation, dispatch and energy operations. Diagrams define candidate boundaries and contracts. Desk review is complete; implementation and validation remain pending.

1.1 Logical containers, request paths and event distribution



Dashed route: I05 equipment feedback via C05 adapters. Solid arrows: principal requests / committed events.

1.2 Runtime deployment, isolation and recovery



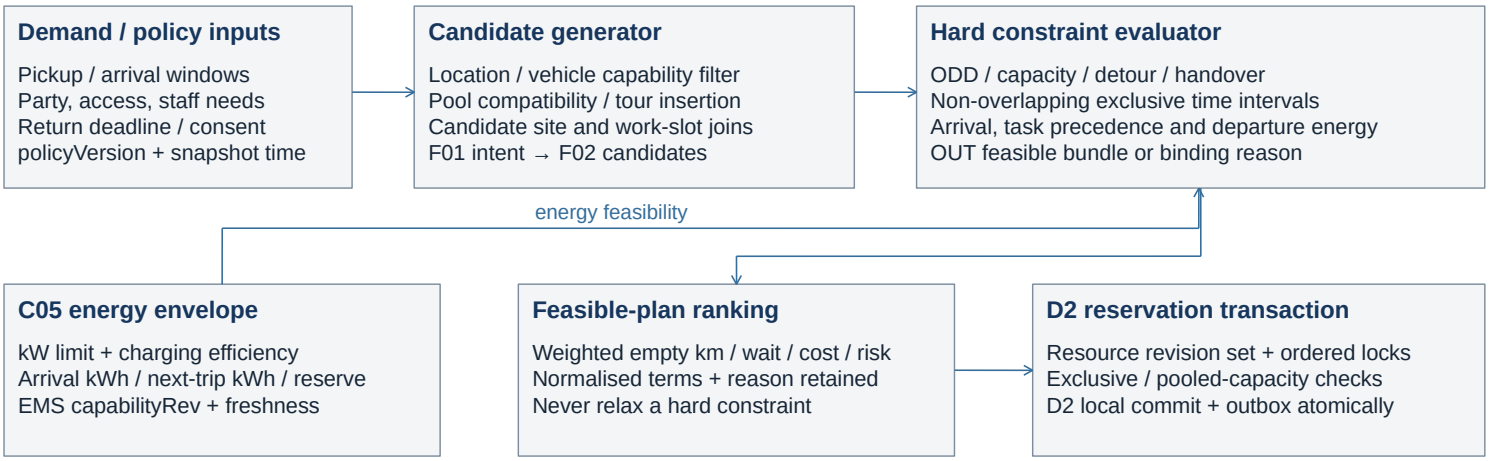
Connector failure may leave local intake available; confirmation still requires all mandatory dependencies. Stale vehicle/site state blocks affected capacity. Replay must use separate projections and idempotent effect handlers.

G2 decisions: hosting zones, failover mechanism, protocols, RPO/RTO, latency and freshness budgets. The primary service integrates booking, dispatch and energy. Equipment/JV/eVTOL options attach through capability interfaces.

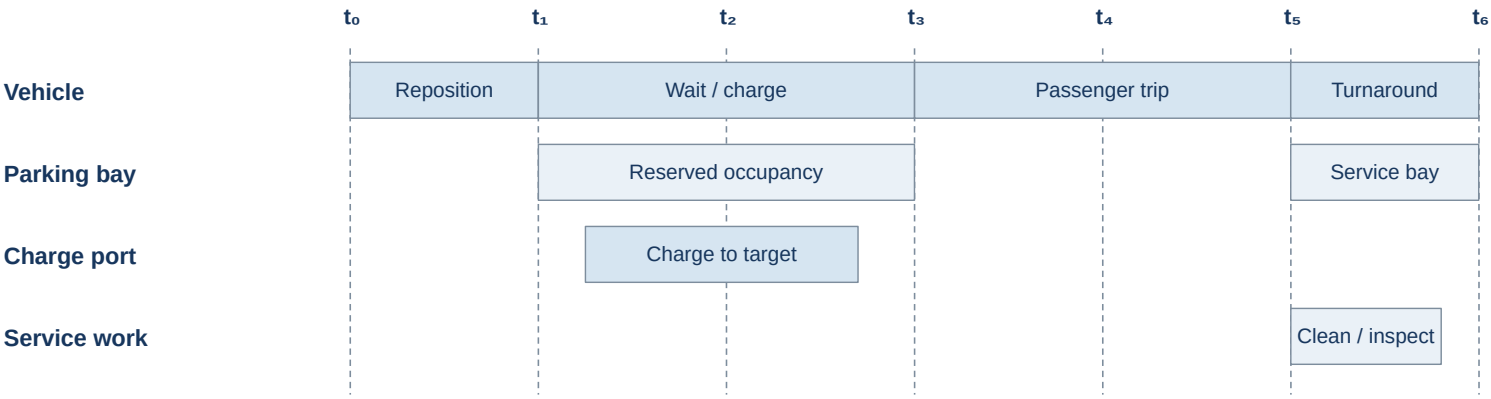
Joint scheduling, resource timing and states

Reservation feasibility includes the complete cycle: pre-position, waiting, passenger service, cleaning or inspection, energy replenishment and the next departure.

2.1 C03 component decomposition and constraint propagation

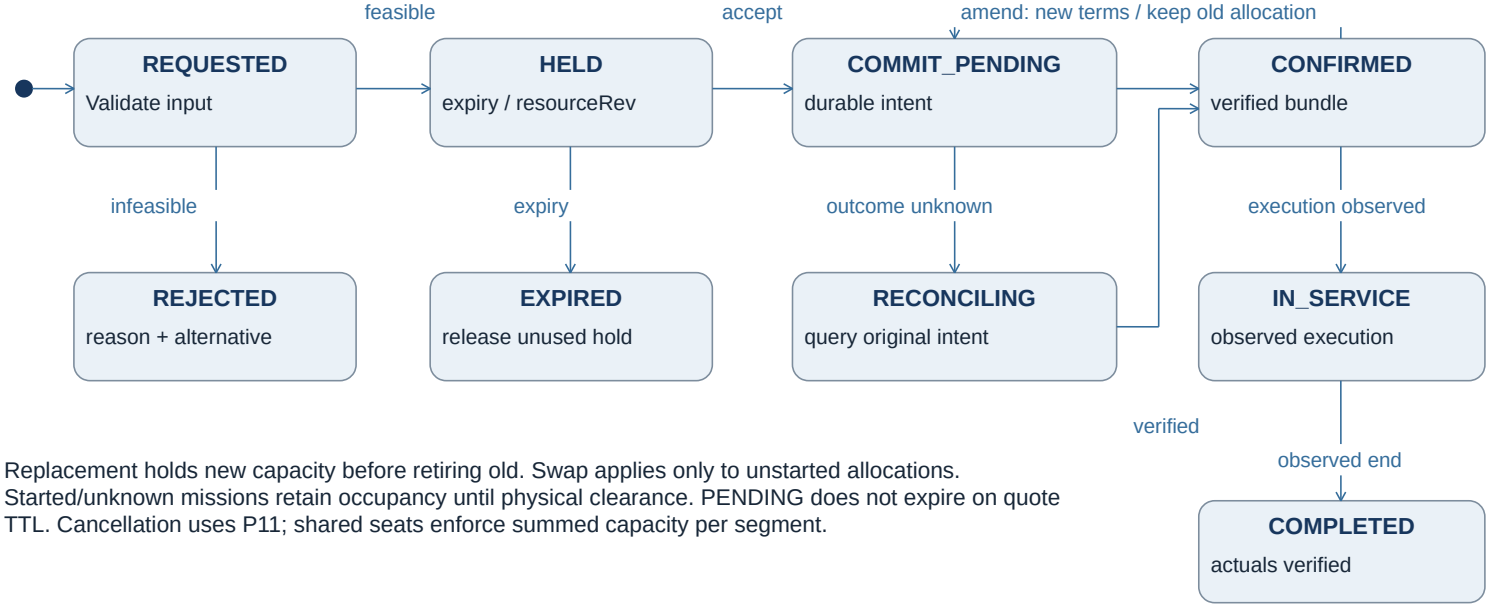


2.2 Coordinated timeline: one vehicle, one bay and one charging session



Departure guard: all required tasks finish before departure; precedence permits only approved overlap. Electrical branch: E_dep = E_arr + Σ(η_ch P_AC Δt_h) − E_aux, with E_trip + E_reserve ≤ E_dep ≤ E_usable,max. kW × hours = kWh.

2.3 Selected reservation transitions and replacement constraints



Replacement holds new capacity before retiring old. Swap applies only to unstarted allocations. Started/unknown missions retain occupancy until physical clearance. PENDING does not expire on quote TTL. Cancellation uses P11; shared seats enforce summed capacity per segment.

Fleet capacity includes empty travel, turnaround and replenishment. Evaluate each zone and time interval against reachable supply. Reserve hub resources separately: bays, connectors, power, work slots and staff. Capacity and coverage checks: P25–P26.

Entity relationships and message contracts

Identifiers, cardinalities, revisions and uniqueness constraints preserve the chain from an accepted journey to resource use and settlement. Example payloads define proposed contracts.

3.1 Domain model with keys and ownership

PlanReference / C01

PK tenantId, planRefId
provider, accountRef, sourceEventId
instanceId, sourceRevision
consentId, purpose, expiresAt
startUTC, endUTC, IANA zone
Unique within tenant:
provider/account/event/instance

0..1 : N

Booking / C02

PK tenantId, bookingId
FK planRefId (optional for app input)
bookingRev, termsRev, state
riderRef, pickupWindow, destination
currentBundleId, intentId
CAS: expectedBookingRev

1 : N

ResourceBundle / C03

PK tenantId, bundleId
FK bookingId, replacesBundleId?
bundleRev, state, expiresAt
policyVersion, inputSnapshotRef
intentId, idempotencyKey, result
Unique: tenant/intent/key

1 : N

1 : N

ActivitySaga / WorkOrder

PK tenantId + sagaId / workOrderId
FK bookingId, bundleId?
providerRef, slotRef, staffRefs[]
conditionsRev, acceptedTermsRev
state, owner, dueAt, evidenceRef
Tasks: cleaning / inspection / handover

ResourceInterval / C03

PK tenantId, intervalId
FK bundleId, resourceId
startUTC, endUTC [half-open]
capacityType, quantity, unit
reservationRev, allocationState
Exclusive: no overlap; pooled: sum ≤ capacity

N : 1

Resource / Capability

PK tenantId, resourceId
kind: vehicle/bay/charger/staff
siteId?, operatorRef, contractRef
capabilityRev, status, observedAt
capacity, constraints, quality
Availability needs current evidence

Bundle references below: 1:N Mission commands • 1:N EnergyPlan revisions • 1:N verified usage and ledger entries

Mission / Command / C04

PK tenantId, commandId
FK bundleId, vehicleId
missionRev, routeRef, validUntil
deliveryState, executionState
statusSeq, observedAt, receivedAt
Same commandId: one immutable payloadHash

EnergyPlan / C05

PK tenantId, planId, planRev
FK bundleId, siteId
capabilityRev, powerIntervals[]
departureMinKWh, expiresAt
acceptedRev, executionEvidence
Plan revision ACK ≠ delivered energy

UsageFact / Ledger / C07

PK tenantId + usageId / ledgerEntryId
FK bundleId, contractRev
meterPointId, interval, qty, unit
quality, sourceRev, payerRef
correctionOf?, billRev, currency
No destructive overwrite of actuals

3.2 Command, result and event examples

I03 CommitBundle / command

```
{
  tenantId: "operator-A",
  intentId: "amend-41",
  bookingId: "bk-18",
  expectedBookingRev: 6,
  newBundleId: "rb-52",
  oldBundleId: "rb-37",
  expectedResourceRevs: [104, 28],
  acceptedTermsRev: 3,
  capabilityRev: "cap-9",
  idempotencyKey: "amend-41"
}
```

Same key + different payload:
reject IDEMPOTENCY_CONFLICT.

I03 Result / guarded outcome

```
{
  intentId: "amend-41",
  state: "COMMITTED",
  bundleId: "rb-52",
  bundleRev: 105,
  intentIdForD1: "amend-41",
  committedAt: "<UTC instant>",
  externalReleaseJobs: ["j-4"]
}
```

Other explicit outcomes:
HELD / REJECTED / PENDING
reasonCode + retryability
Unknown ≠ rejection.

Event envelope / all domains

```
{
  eventId: "evt-902",
  type: "AllocationCommitted",
  schemaVersion: 1,
  tenantId: "operator-A",
  aggregateId: "rb-52",
  aggregateRevision: 105,
  occurredAt: "<UTC instant>",
  receivedAt: "<UTC instant>",
  correlationId: "amend-41"
}
```

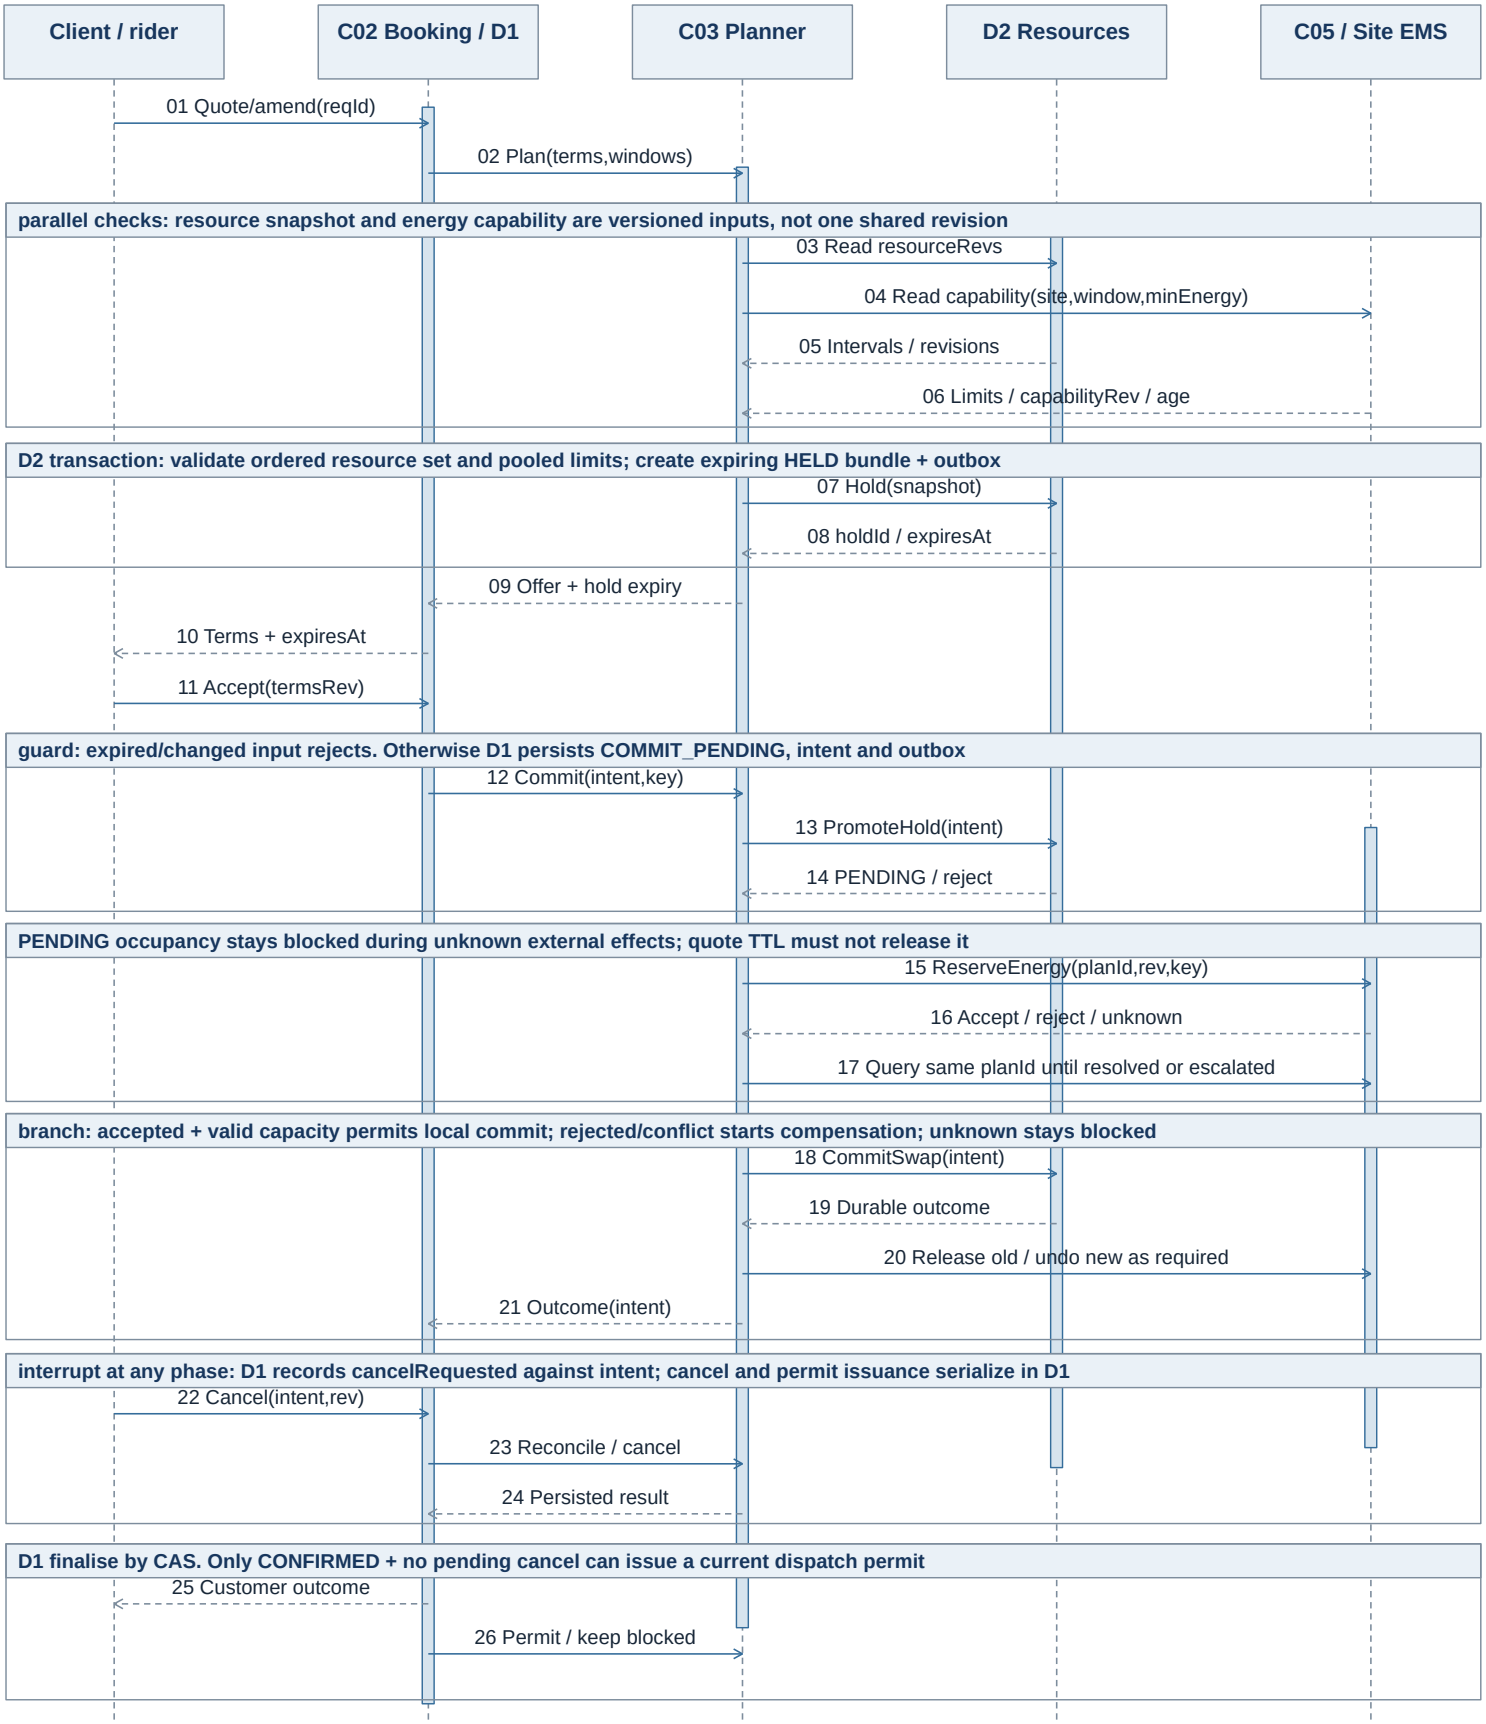
Inbox dedup → revision check
→ apply / fetch gap / quarantine

Keys/FKs include tenant scope; validate references within that scope. expectedResourceRevs maps to the ordered resource set. D1 owns bookingRev. Atomic guarantees cover D2 only. Examples are conceptual payloads, not complete API schemas.

Booking saga, cancellation and dispatch gate

C02 owns booking intent and cancellation. D2 makes only local atomic allocation changes. External EMS effects use a reconciled saga; no distributed atomic commit is assumed.

4.1 Hold promotion, external outcome and authoritative dispatch permit

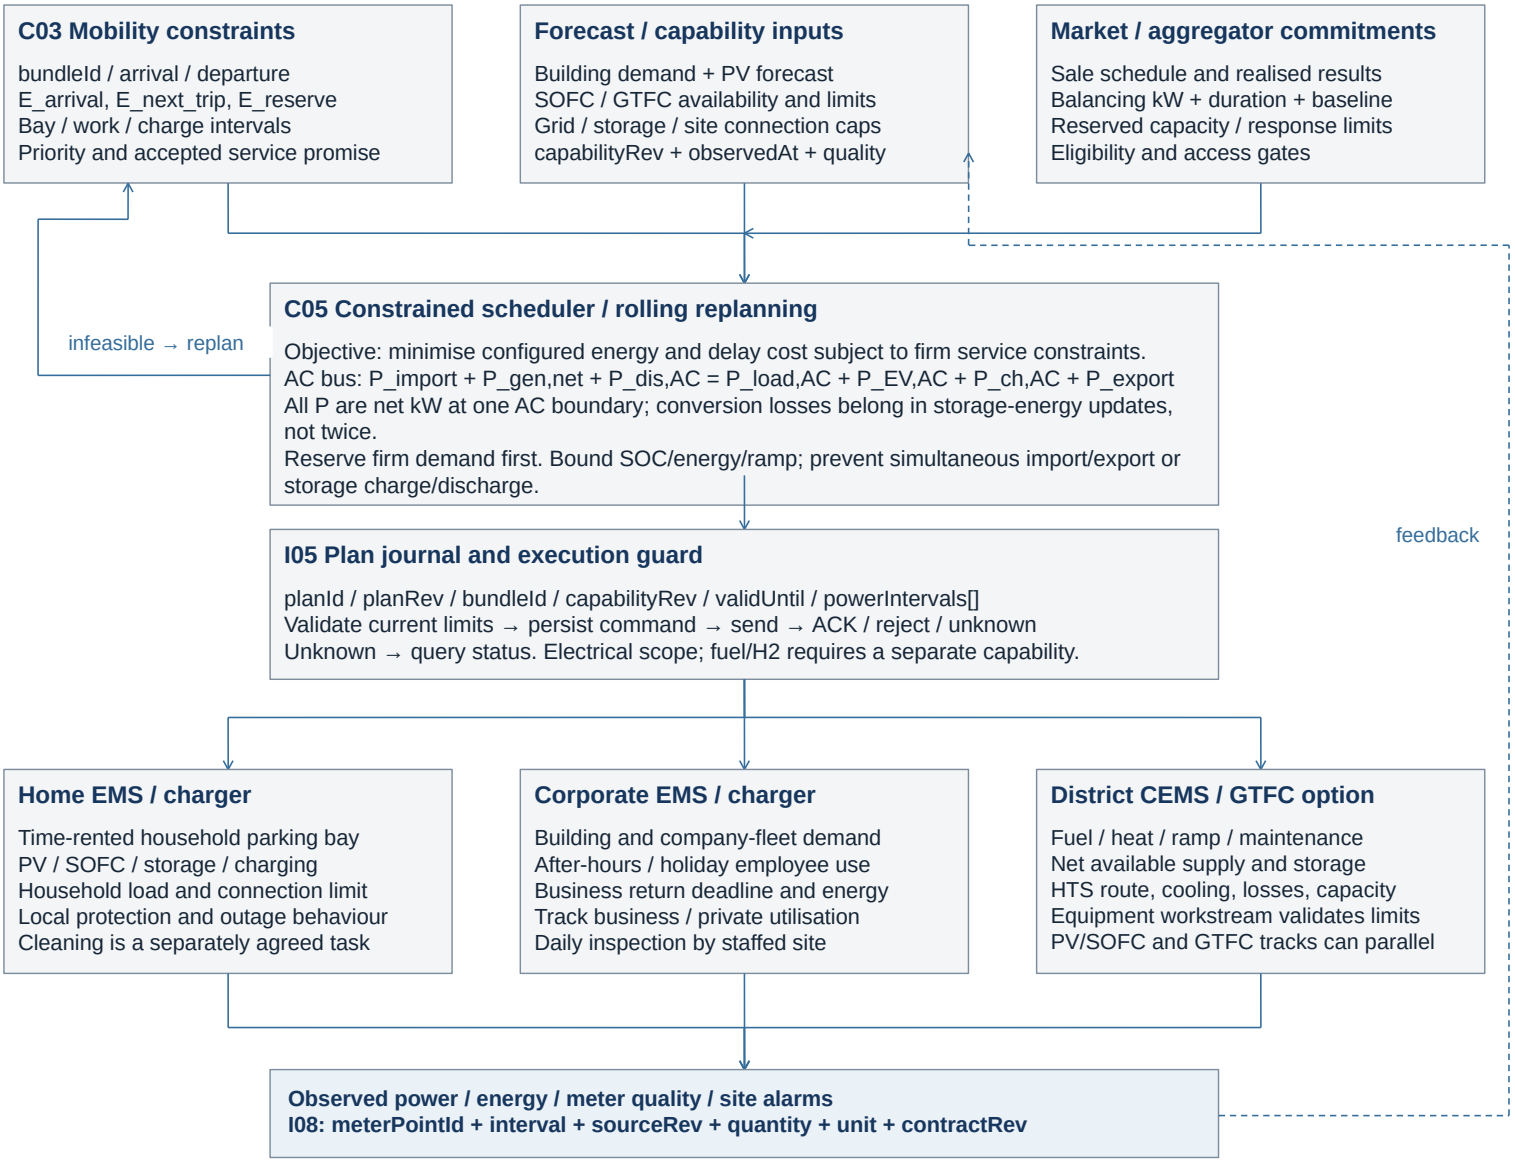


C04 obtains the current D1 permit before releasing a mission. Cancellation after permit issuance becomes cancel/replan of an in-flight mission, not a guarantee of zero movement. Release old physical capacity only after vehicle/EMS clearance. Failed compensation stays owned and reconciled.

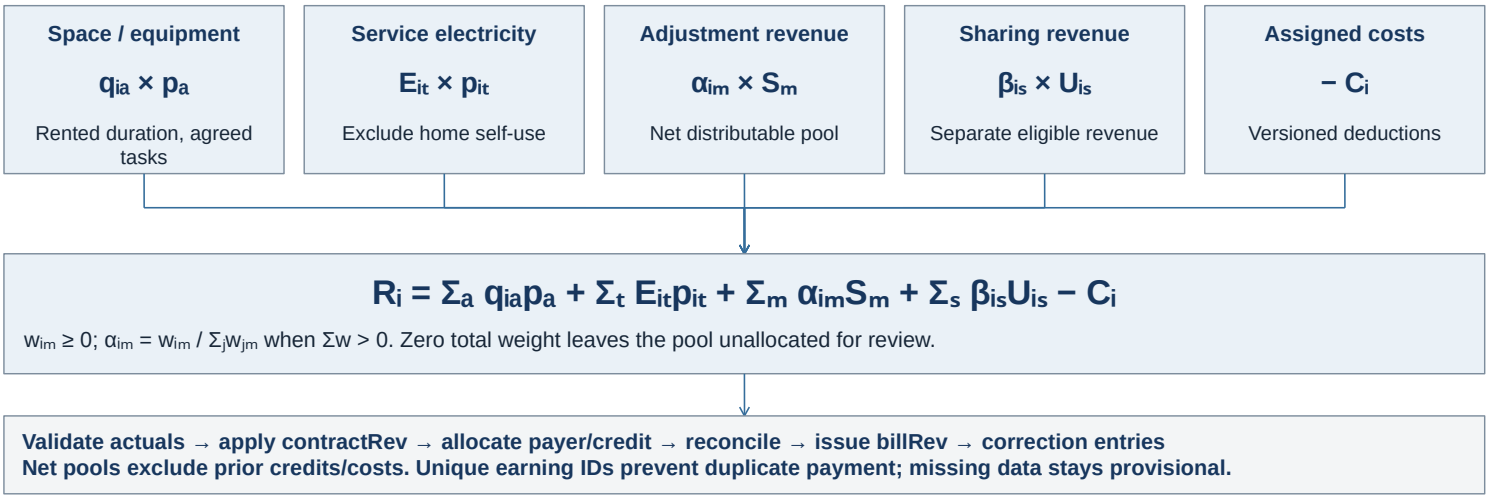
Energy control loops and contribution accounting

The mobility schedule imposes departure-energy constraints. C05 allocates feasible power, reconciles equipment response and sends shortages back to C03. C07 settles measured contributions.

5.1 Hierarchical control and measurement paths



5.2 Settlement dataflow and household contribution formula

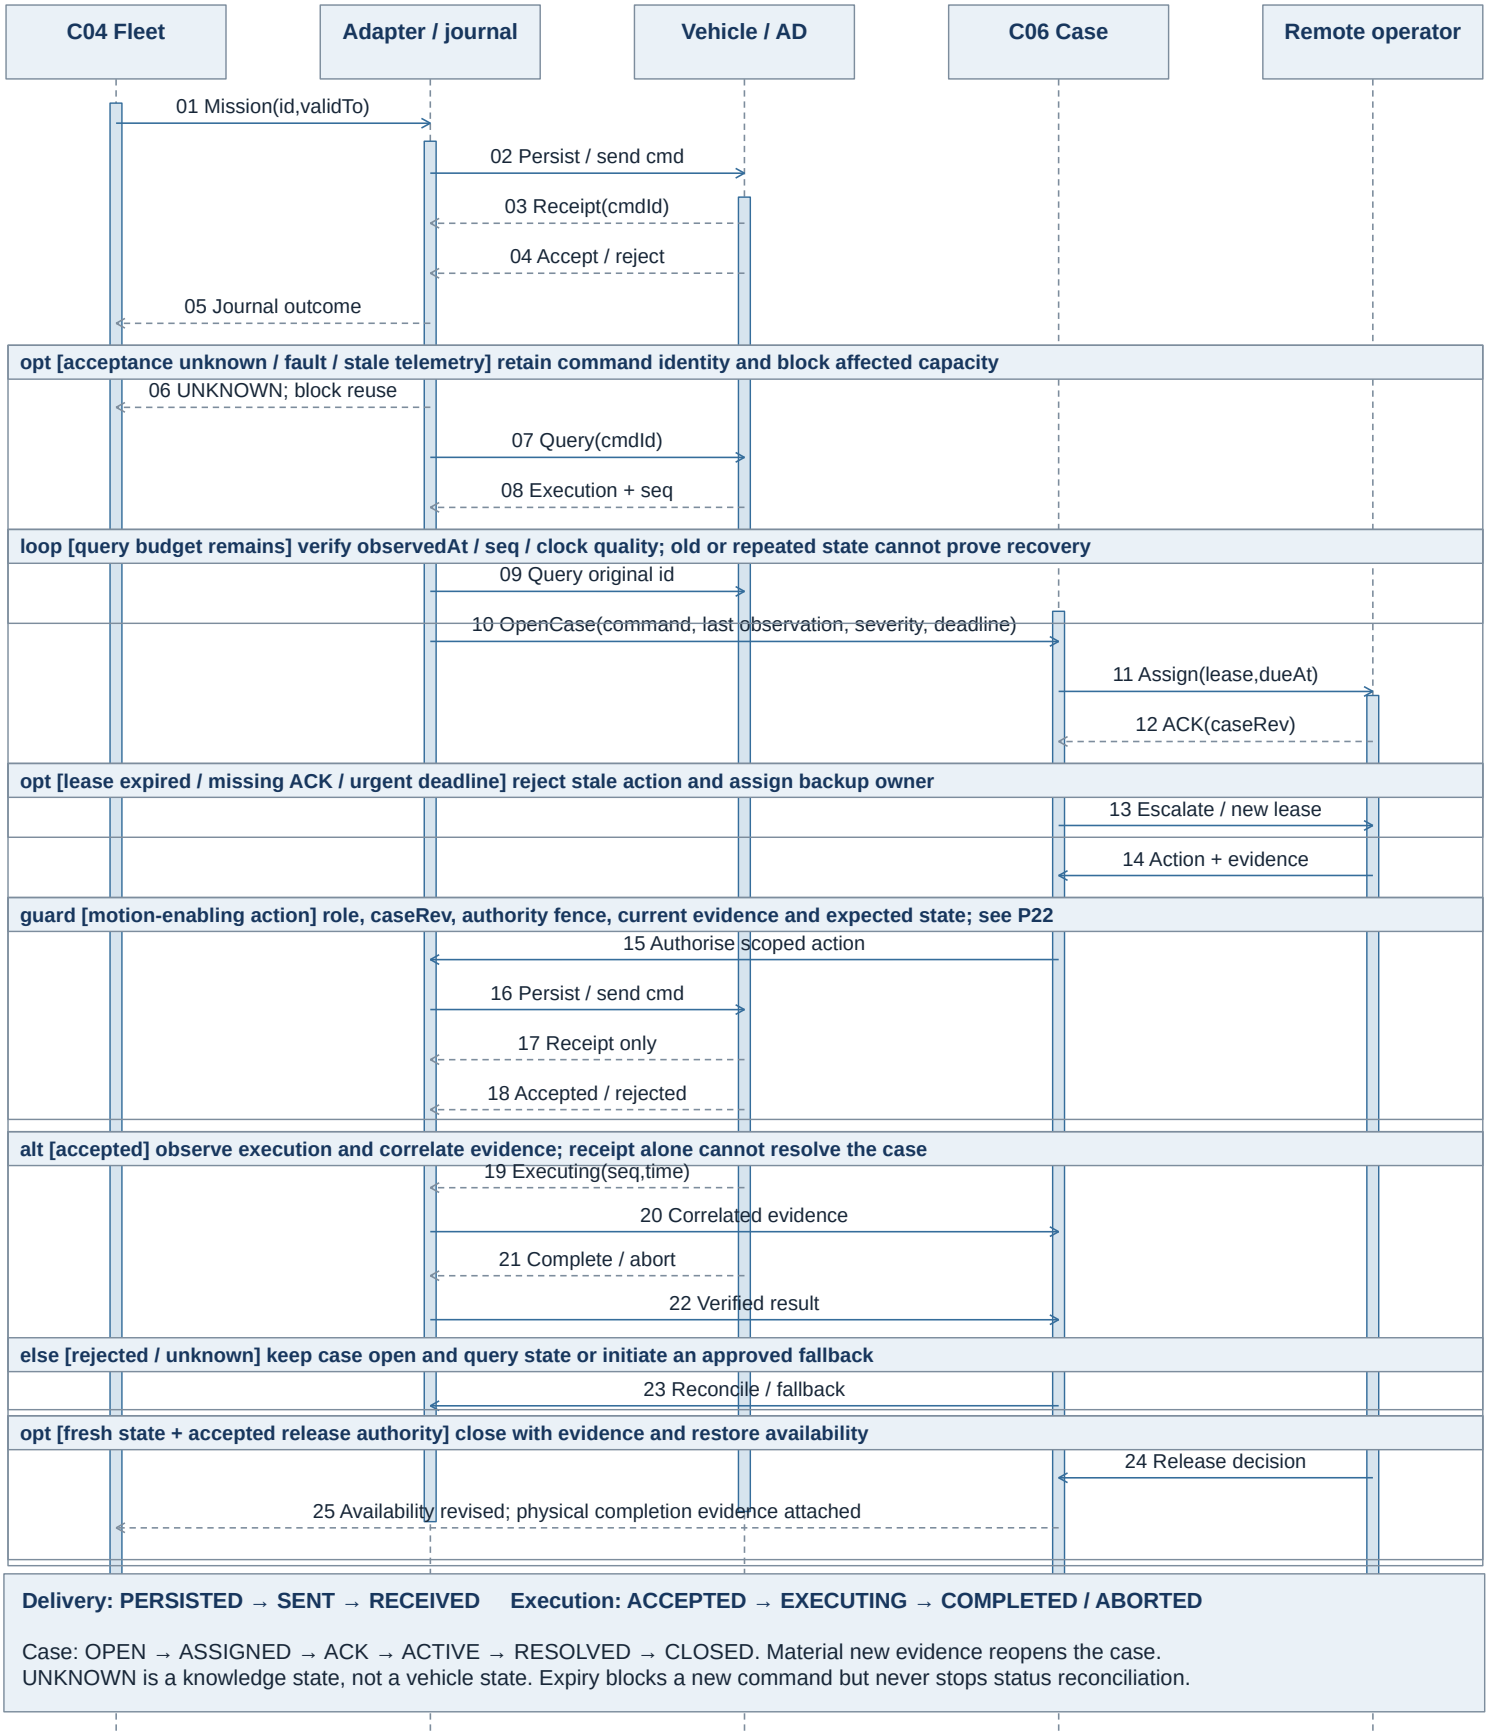


Replenishment queues, site derating and reserve energy constrain each next departure (P26). Settle measured delivery under a versioned contract. Keep net pools, credits and assigned costs non-overlapping. Availability kW and delivered kWh remain distinct.

Remote mission protocol and exception handling

I04 delivery acknowledgment, vehicle execution and I07 case resolution carry different evidence. Permission checks and local vehicle authority constrain every remote action.

6.1 Persistent command, uncertain execution, case ownership and recovery



Delivery: PERSISTED → SENT → RECEIVED **Execution: ACCEPTED → EXECUTING → COMPLETED / ABORTED**

Case: OPEN → ASSIGNED → ACK → ACTIVE → RESOLVED → CLOSED. Material new evidence reopens the case.
UNKNOWN is a knowledge state, not a vehicle state. Expiry blocks a new command but never stops status reconciliation.

Vehicle/ADS owns driving and fallback. Read-only queries and emergency escalation remain available with stale data. Case closure follows P18; motion requests use P22 action-specific guards. Rejected/unknown effects stay reconciled.

Remote architecture and control authority

RM-01 / RM-02. Level 4 ADS retains DDT and fallback authority. Remote monitoring, contextual assistance and passenger support use separate permissions and evidence paths.

7.1 Onboard functions, remote channels and command boundary

VEHICLE TRUST ZONE / device identity, boot epoch and capability revision

ADS / local safety authority

OEDR: detect and respond to road events
DDT: tactical / operational driving
Perception, localisation, motion control
DDT fallback initiates local response
MRM seeks MRC with achieved-state evidence. Zero speed alone is not MRC.

Edge adapter / durable journal

Vehicle state, ODD and diagnostic feed
Command identity, expiry and dedup
Receipt / acceptance / execution split
Safety checks remain onboard
Prioritised uplink + event buffer
No cloud override of ADS constraints

Media / passenger interface

Exterior views and calibration IDs
Scene / frame capture timestamps
Optional approved cabin channels
Passenger SOS / accessible voice
Independent media-health evidence
No stale frame rendered as live

Telemetry / alarm channel

Reliable events + heartbeat + seq
Gaps, clock quality, age and validity
A missing signal is explicit uncertainty.

Command / result channel

Authenticated, scoped, expiring action
Current lease + capability + scene
Onboard accept/reject + observed result

Evidence / media channel

Adaptive video / keyframe evidence
Separate channel health and age
Replay mode disables live actions.

C04 vehicle adapter / I04

Validate schema, identity, unit and seq
Project state with observation quality
Journal command effects and unknowns
Preserve raw events and evidence gaps
Query original IDs after lost responses
Never default missing state to healthy

C06 cases / I07 policy

Correlate vehicle + I09 roadside evidence
Qualified seat + single action lease
Bounded RA advice / escalation
Preserve source quality and conflicts
Fresh vehicle evidence gates release
Service AI cannot authorise motion

Operator / PaaS consumers

Focus scene, history and evidence age
C03 holds affected vehicle capacity
C02 revises journey / passenger offer
C05 revises feasible energy / site slots
No automatic re-entry on reconnect
No release merely because case closed

Roadside sensors / RSU / edge → C10 I09 validation → C06 monitoring evidence [details P15–P16]

7.2 Operational roles and failure containment

Remote assistance and remote driving

RA supplies contextual advice that the ADS can reject after re-evaluating the scene and ODD. Real-time remote DDT is outside this baseline. Adding remote driving requires a distinct operating envelope, control arbitration, communications safety analysis and operator / approval evidence.

Operational and legal responsibility

Map system roles to the permitted operation, designated supervisor and field-response organisation. Passenger support cannot override door interlocks or authorise motion. Vehicle-specific link-loss and fallback behaviour remains local and does not depend on a remote reply.

Active authority

Durable cases / command journal
Replicated evidence + backlog
Vehicle fence ACK gates new actions

Standby promotion

Acquire higher authority epoch
Require vehicle fence ACK (P20)
Reconcile prior effects before action

Degraded service

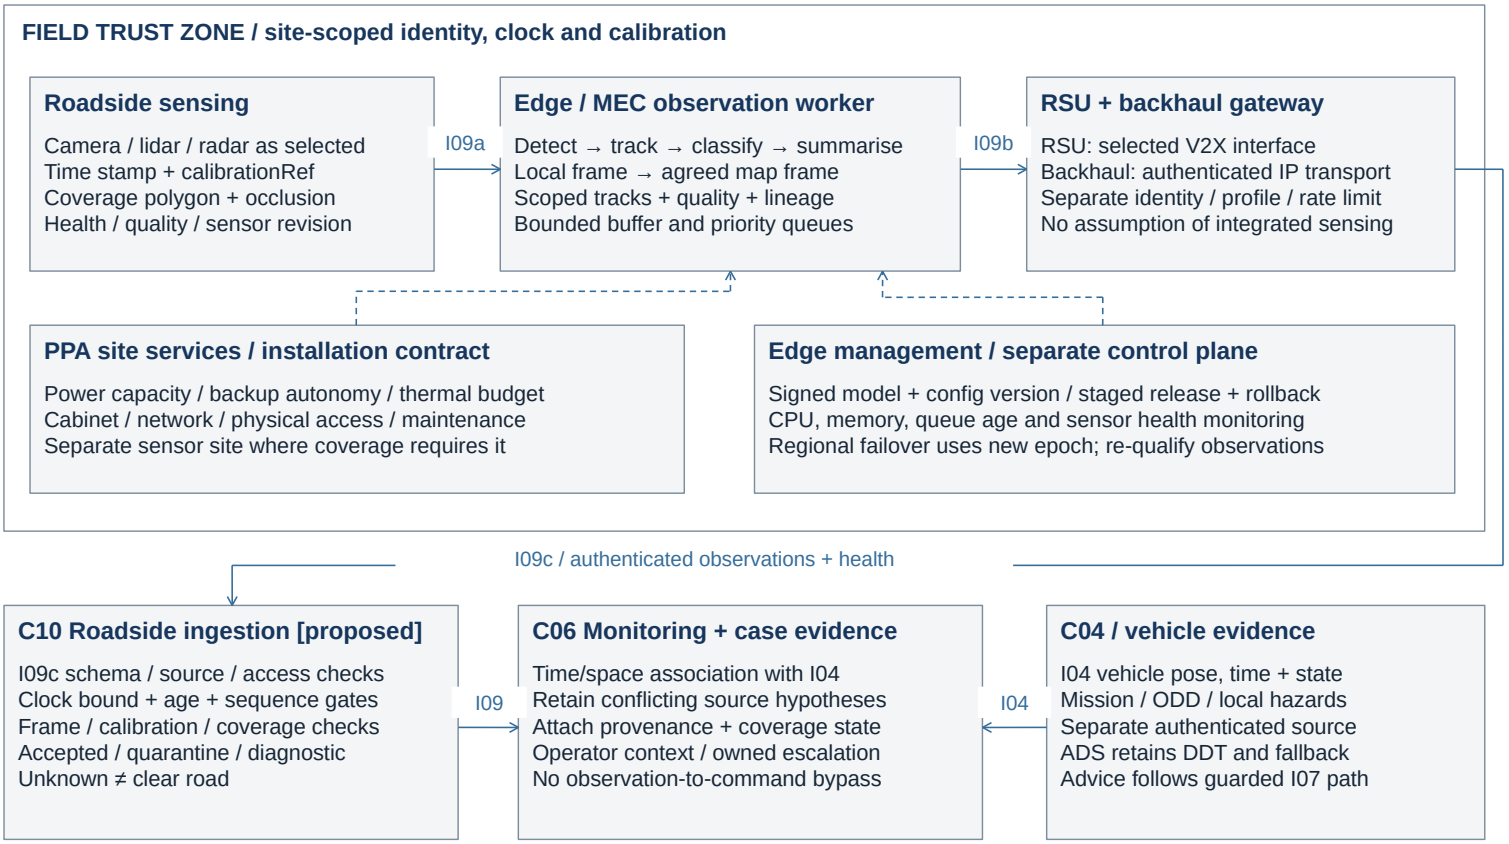
Hold affected assignments
Preserve priority safety traffic
Escalate without hiding uncertainty

Acceptance: independently partition telemetry, media, primary centre and workstation. Old-epoch requests reject after the vehicle installs its fence. Test durable deduplication and unknown outcomes. Reconnect and case closure require the release predicate on P18 before availability can change.

Roadside sensing and edge computing architecture

RM-13. Roadside observations supplement vehicle evidence. Sensor, compute and communication functions have separate ownership. Proposed I09 feeds monitoring and operational review, not direct actuation.

8.1 Deployment topology and data / management plane separation



8.2 Proposed observation contract and association keys

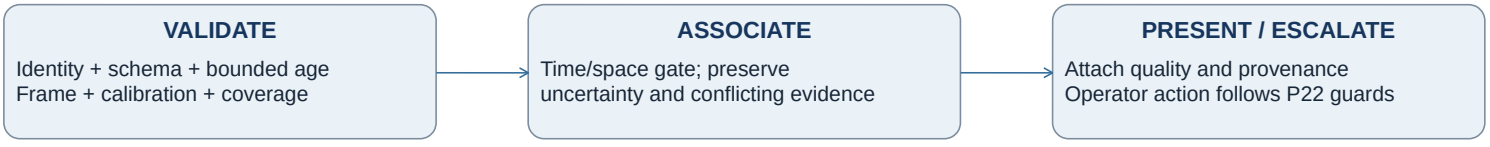
ObservationEnvelope / I09c

tenantId, siteId, sourceId, sensorIds[], bootEpoch, seq
observedAt, emittedAt, clockErrorBound_s, expiresAt
trackKey = sourceId / bootEpoch / localTrackId
frameId, mapRev, calibrationRev, coverageRef
position_m, velocity_mps, uncertainty / covariance
class + confidence, visibility, modelRev, evidenceRef
Signature / authenticated transport + payloadHash

CoverageHealth + association evidence

coverageId / polygon / exclusions / sampledAt
Sensor availability, clock validity, calibration status
Occlusion / weather / queueAge / processing status
Vehicle association: spatial-temporal gate + uncertainty
Candidate link has confidence, method and expiry
No global identity assumed from local track IDs
Multiple reports retain source lineage to avoid double count

8.3 Evidence eligibility, overload and storage boundaries



Invalid / stale / overloaded

Quarantine observations, publish coverage degradation.
Reserve queue capacity for health / critical events.
Drop policy is explicit; historical replay cannot act live.

Evidence / privacy / regional placement

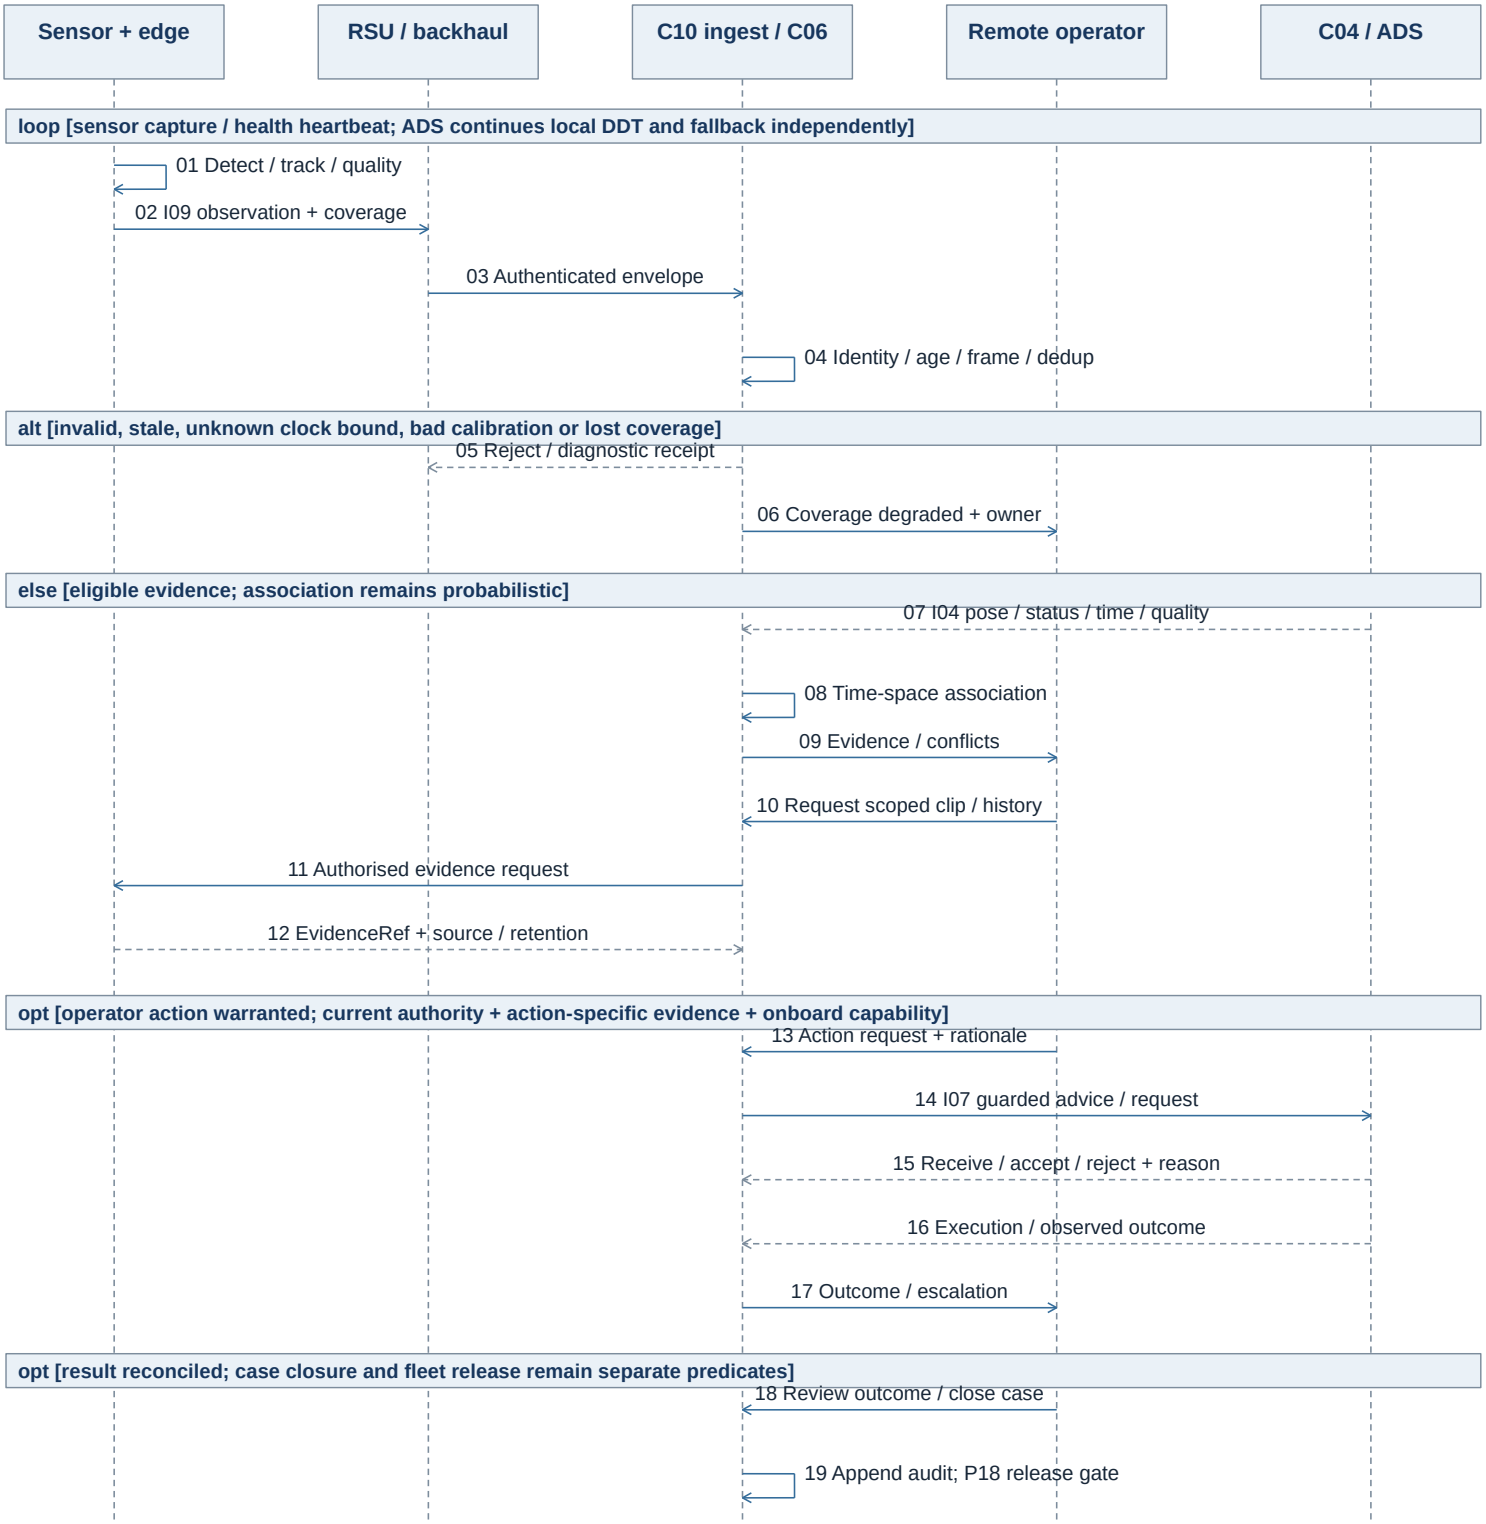
Raw clips by purpose + access + retention contract.
Edge summaries first; authorised evidence on request.
Regional workers scale independently of ADS control.

Thresholds for age, association, coverage and queue size are agreed per use case. CPM mapping and radio profile require a separately verified adapter. Proposed monitoring path does not claim collision avoidance or perception completeness.

Roadside observation and remote monitoring sequence

RM-14. A roadside event becomes scoped monitoring evidence only after validation. The sequence retains independent ADS behaviour, conflicting evidence and unknown outcomes.

9.1 Observation, evidence correlation and guarded response



9.2 Fault branches and planned verification

LINK / EDGE LOSS
Heartbeat expiry → coverage UNKNOWN; retain source time.
No stale frame displayed as current. Reconnect uses new epoch, resync + health checks; replay is historical only.
Test disconnect, reboot, reordering, duplicate delivery.

CONFLICT / OVERLOAD
Keep both source claims; no silent override of ADS state.
Overload reports queue age and dropped observations.
Test occlusion, weather, false association and saturation.
Measure quality, end-to-end age and operator workload.

Telemetry coverage and alarm decision logic

RM-03 / RM-04. Capability-specific signals feed versioned detection and triage. Alarm priorities are operational response classes, not ASIL classifications. Every action uses current, valid evidence.

10.1 Signal families, decision context and missing-data behaviour

ADS / ODD	adsMode, engagementState, oddVerdict/reason, oddProfileRev Route / geofence / weather / illumination constraints as supported.	ODD or mode inconsistency opens a case. Cloud context never widens ODD.
Motion / localisation	pose + coordinateFrame, speed, heading, localisationHealth Covariance and integrity bounds remain distinct supplier quantities.	Detect unexpected stop / deviation. Uncertain pose cannot prove location.
Perception / planning	sensorHealth, occlusionFlags, plannerStatus, blockageReason assistanceRequest + sceneRev + mapRev + frame capture time	Review ambiguity in context. Generic “AI confidence” cannot permit motion.
Chassis / fallback	brake / steering / propulsion health, emergencyStopStatus fallbackReason, mrmState, mrcAchieved, standstill / secureState	Separate manoeuvre progress from achieved MRC. Failed fallback escalates.
Passenger / mission	missionId/rev, pickup/drop-off state, door/ramp/restraint status SOS, occupancy quality, accessibility and handover exceptions	Local interlocks prevail. Unknown occupancy requires supported confirmation.
Energy / link	SOC, usableEnergy, HV/thermal fault, departure reserve margin heartbeat, RTT, loss, jitter, mediaAge, seqGap, clockQuality	Energy shortage revises feasibility. Link loss marks remote state UNKNOWN.

10.2 Evidence validation, alarm lifecycle and response classes

Validate / detect

Identity, schema, unit, epoch/seq
Age / source quality / clock bound
Predicate + persistence + hysteresis

→

Correlate / assign

Vehicle / cause / incident window
Severity + exposure + dueAt
Qualified seat + ACK + lease

→

Act / verify / clear

Guard action and observe outcome
Stable clear condition + release
Reopen on contradictory evidence

P0	Collision / fire / SOS / failed fallback. Immediate escalation, no debounce.
P1	ODD exit / MRM / critical fault / unknown active state. Hold capacity, urgent response.
P2	Assistance / obstruction / boarding failure. Assess exposure and escalate if worsening.
P3	Noncritical service or diagnostic issue. Schedule work under explicit restrictions.

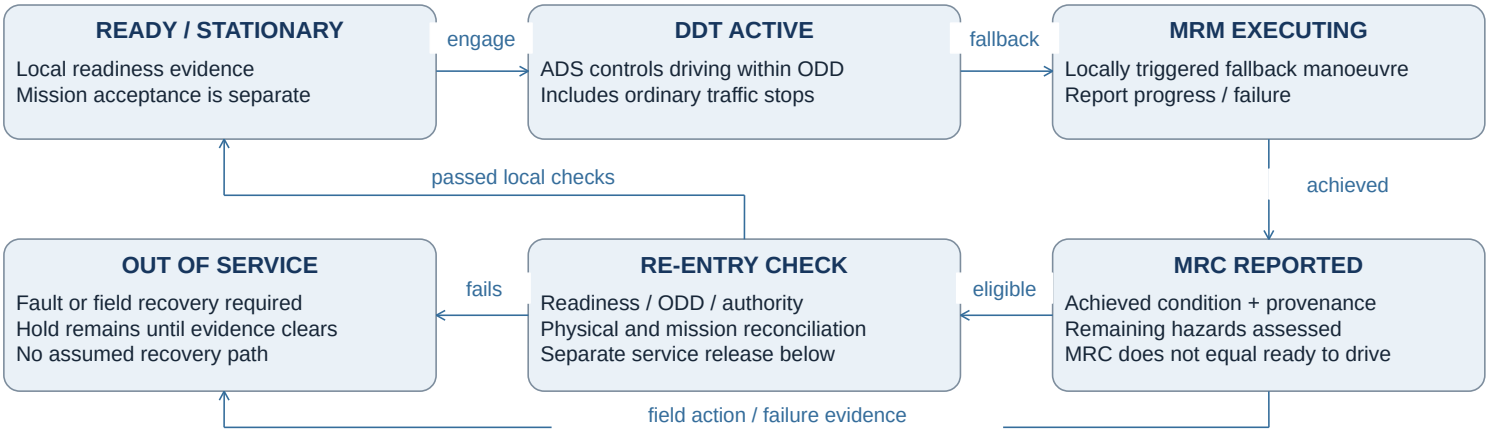
Vehicle I04 / roadside I09 remain separate source envelopes. I09 adds coverageRef, frame/map/calibrationRev, clock bound and source lineage. Preserve original observation time; repeated or replayed samples never refresh it.

Clear only after valid samples span an agreed stable interval. ACK does not clear health. Common-cause grouping retains vehicle exposure and deadlines. Snooze needs reason/expiry and cannot suppress P0. Missing/invalid data raise their own event.

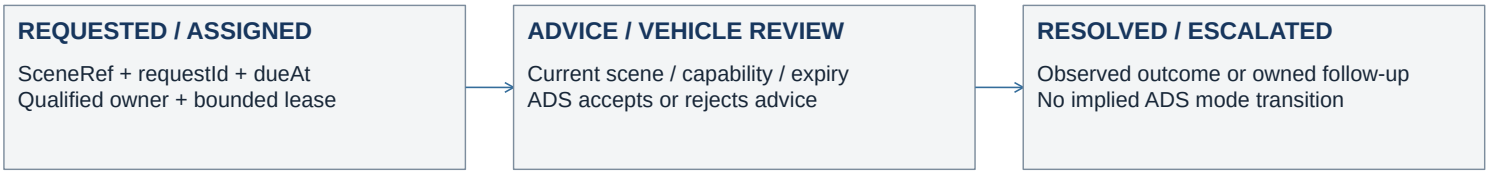
Orthogonal states and guarded return to service

RM-05. ADS operation, assistance workflow, remote knowledge and fleet eligibility are separate dimensions. An ADS can continue its approved local behaviour while a remote assistance request is pending.

11.1 Vehicle-reported operating state: selected transitions



11.2 Assistance lifecycle runs independently of ADS operation



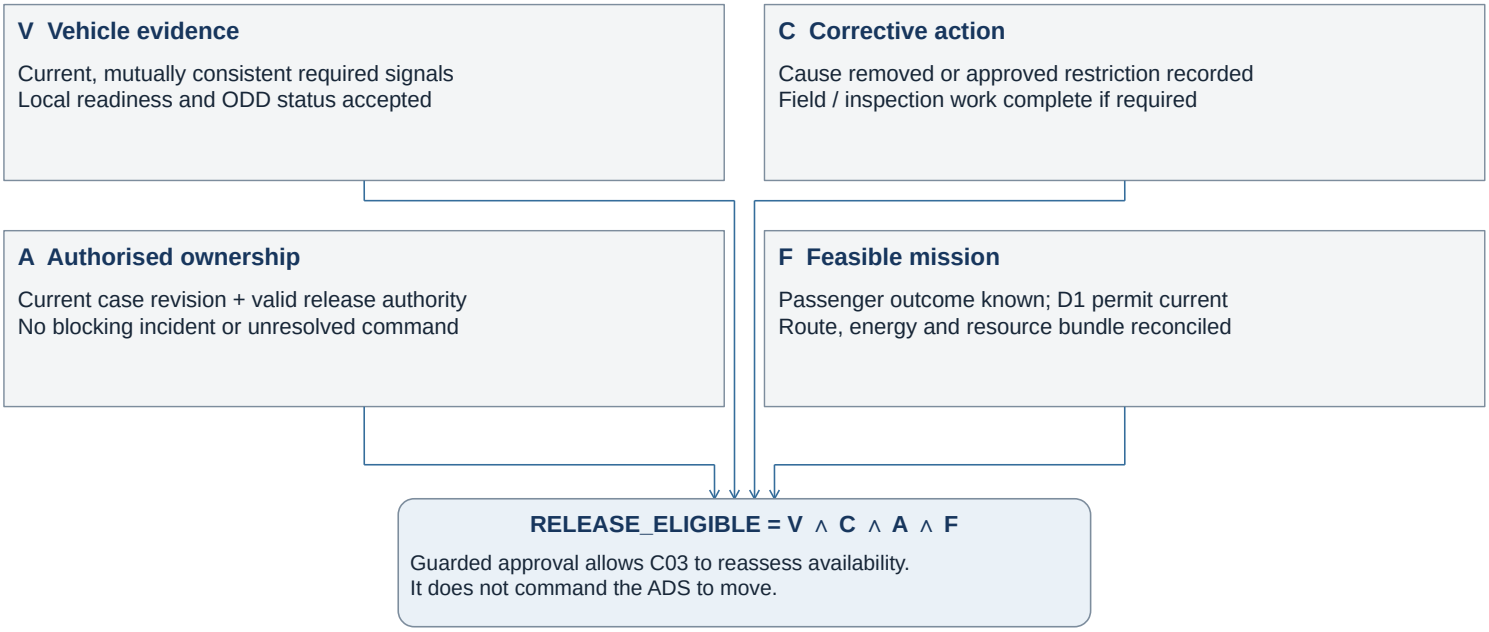
Remote knowledge

CURRENT / DEGRADED / UNKNOWN is per required signal. Fresh data may still report a fault. Unknown physical execution blocks conflicting reuse.

Fleet eligibility / case state

AVAILABLE / RESTRICTED / HELD differs from OPEN / ACTIVE / CLOSED. Closing a case never changes ADS mode or automatically clears an operational hold.

11.3 All release predicates must hold; none substitutes for another

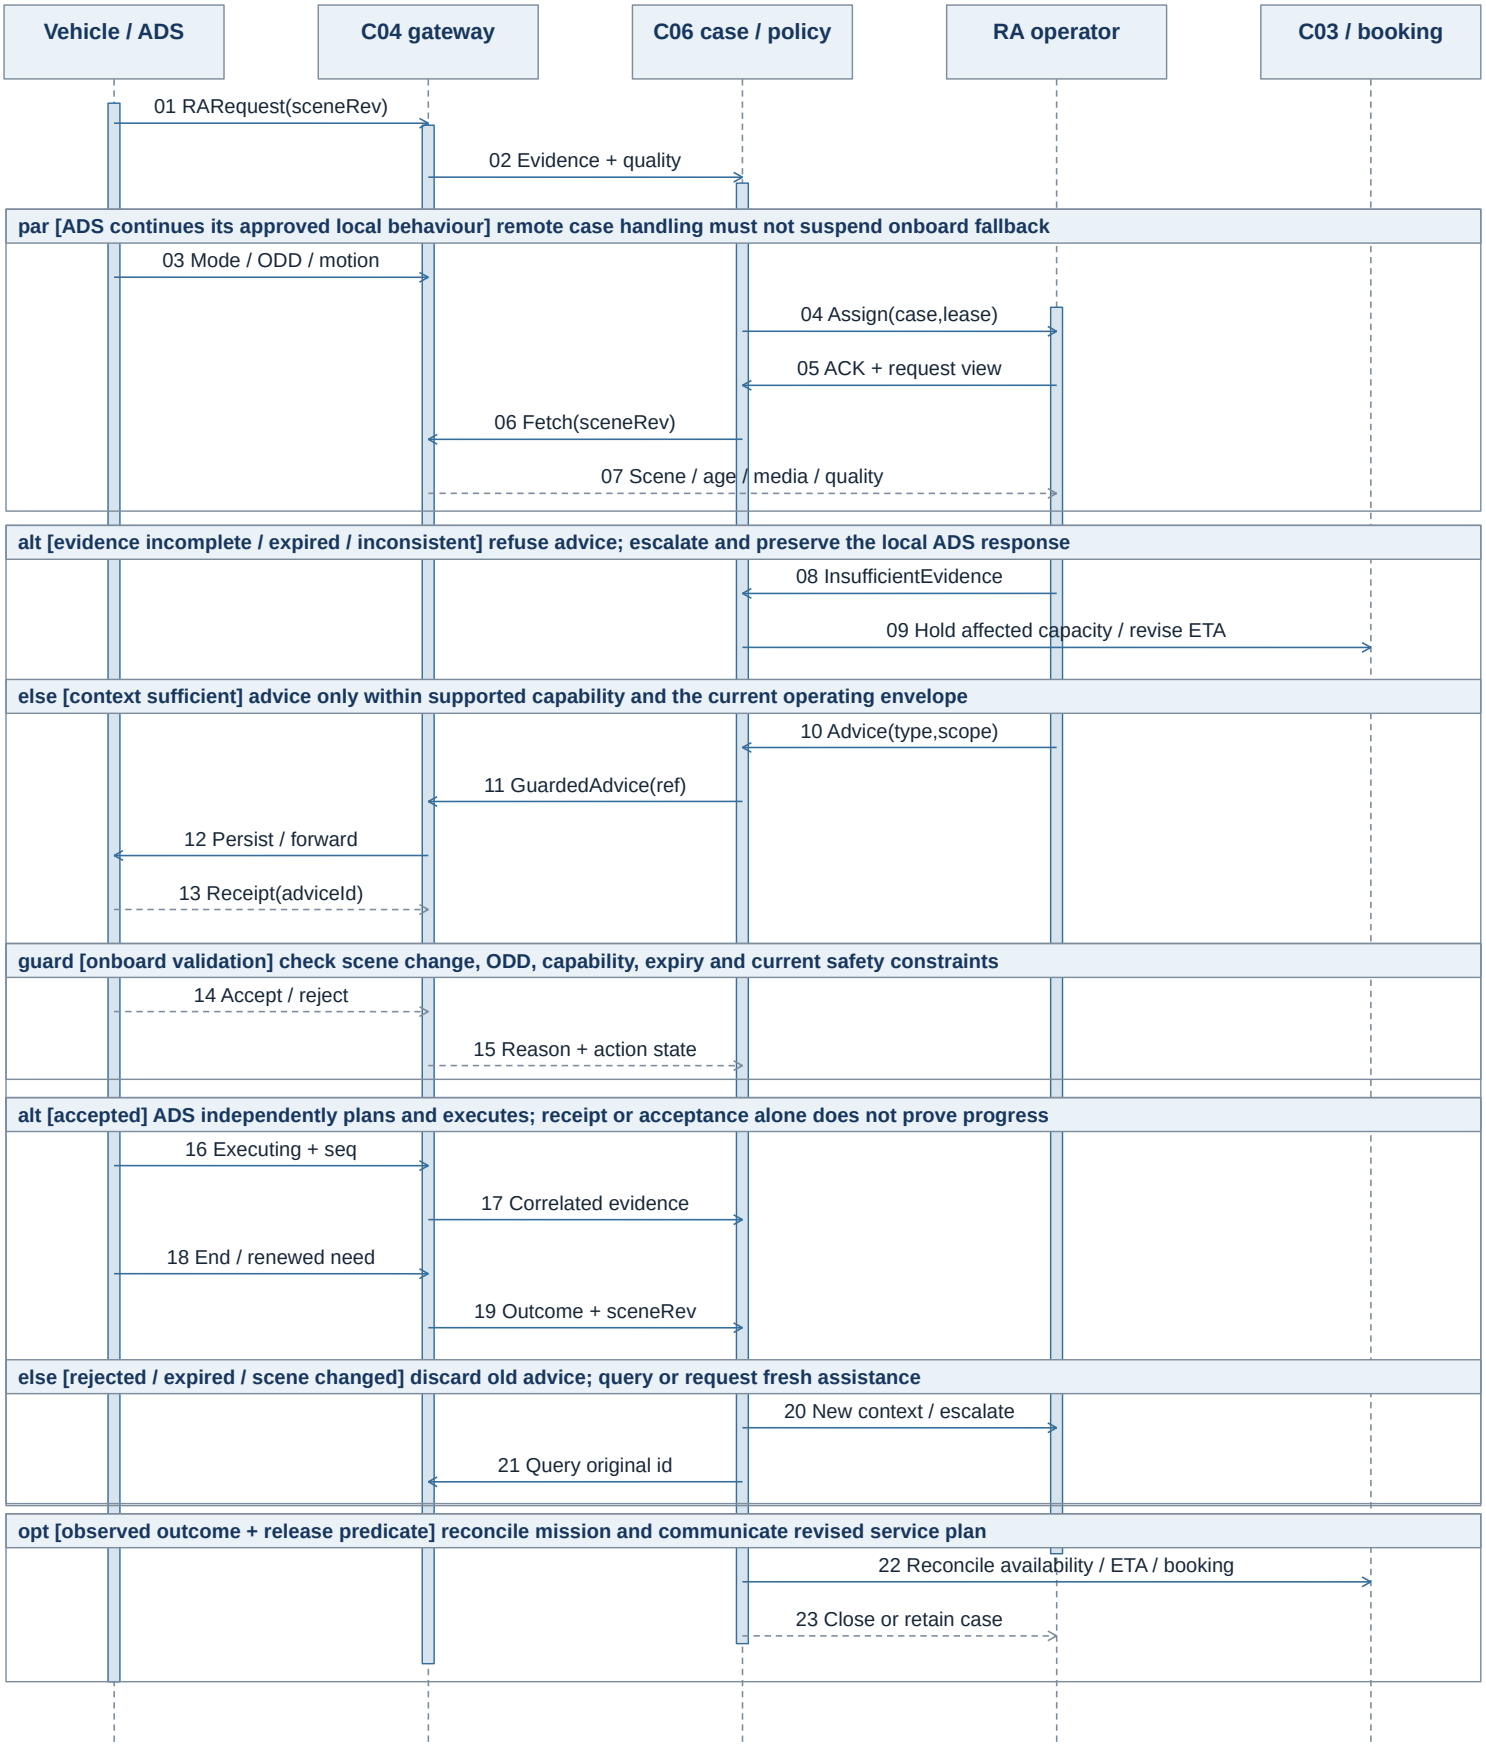


Test cross-products, not only one state chain: DDT_ACTIVE + RA_PENDING; MRC + stale evidence; CLOSED case + critical fault; fresh telemetry + unresolved command. A reported zero speed never suffices for MRC or restart.

Remote assistance: contextual advice sequence

RM-06. Example: the ADS encounters an ambiguous obstruction. A bounded advice transaction carries a scene reference and expiry. The ADS retains trajectory generation, collision checks and motion control.

RM-06A Request, evidence, advice validation and observed outcome

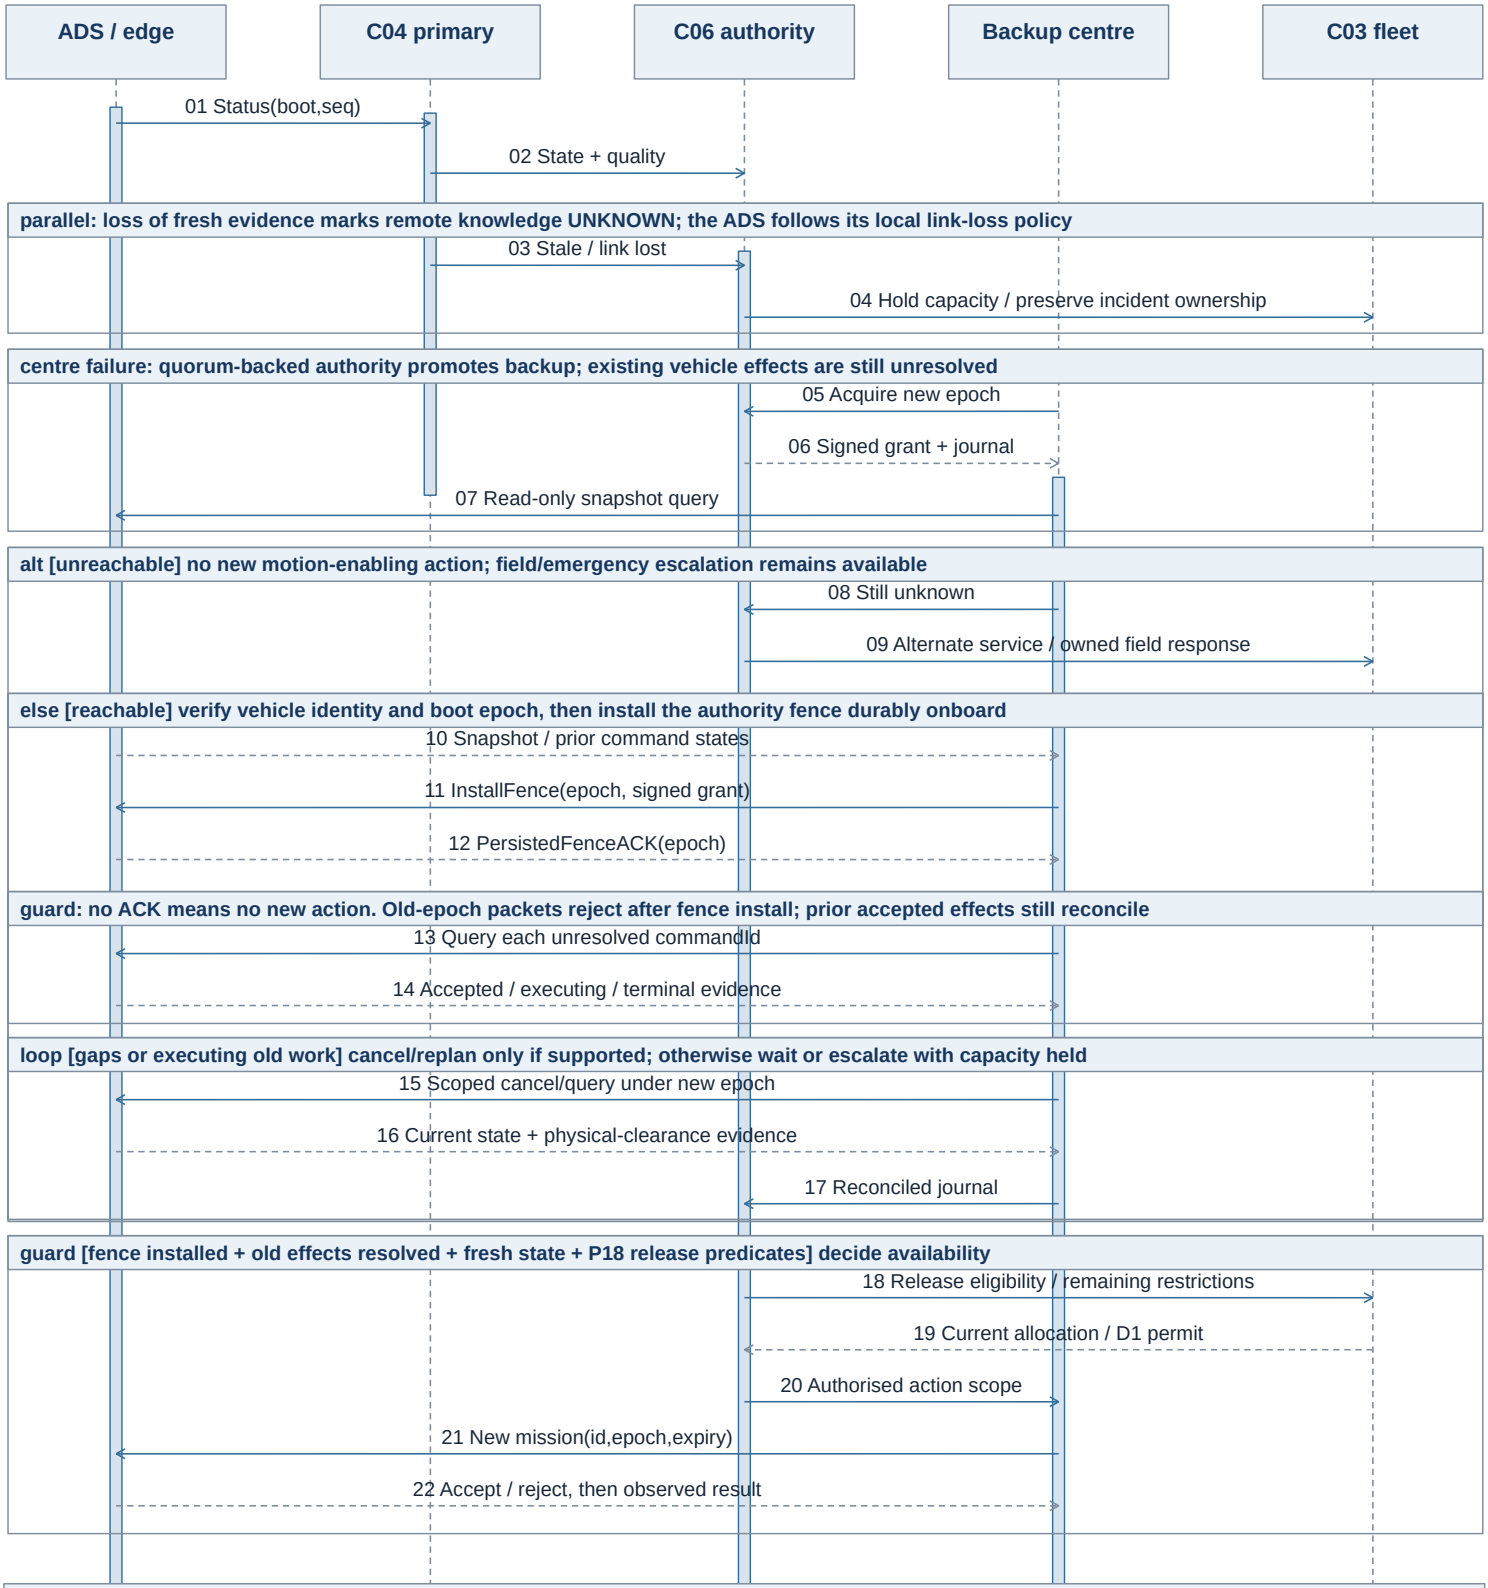


RM-06 invariant: an advice response cannot widen ODD, disable collision avoidance or create a direct actuator path. Timeout/rejection requires reassessment; fresh autonomous resolution can close the case. Executing effects retain their identity.

Link loss, authority fencing and reconciliation

RM-07. Centre authority and vehicle execution require separate evidence. New cloud ownership does not instantly revoke an unreachable vehicle's old authority. Unreachable vehicles remain held.

13.1 Centre change with per-vehicle fence installation



Fence contract: vehicle durably stores its accepted authority epoch and validates signed grants, expiry and command identity. Reboot without recoverable authority/journal evidence blocks new action until resynchronised.
Centre ownership loss blocks new dispatch, but cannot erase an already accepted physical effect. No automatic resend with a new ID. An unsupported fence/dedup capability excludes automated centre handover for that vehicle.

Operator workspace and incident response

RM-08 / RM-09. A focused workspace links fleet triage, ADS evidence, passenger contact and field recovery. Each workstream retains a qualified owner and its own authority.

14.1 HMI wireframe: evidence, authority and operator handover

CENTRE ACTIVE / epoch 42 **Operator O-17 / qualified RA** **Case C-204 / lease current**

Fleet / incident queue

Severity / vehicle / dueAt
Passenger exposure / owner
Oldest unacknowledged case
Common-cause incident link
Seat qualification and health
Escalation always accessible
Colour is not the only cue.

Scene / media / synchronised history

Exterior views / pose / route / scene objects
Map and calibration revision / ADS intended path
LIVE, DELAYED or REPLAY with capture age
Missing view explicit / source quality visible
Align observedAt and receivedAt / mark gaps
Replay disables action. Annotation keeps author, time, evidence reference and case revision.

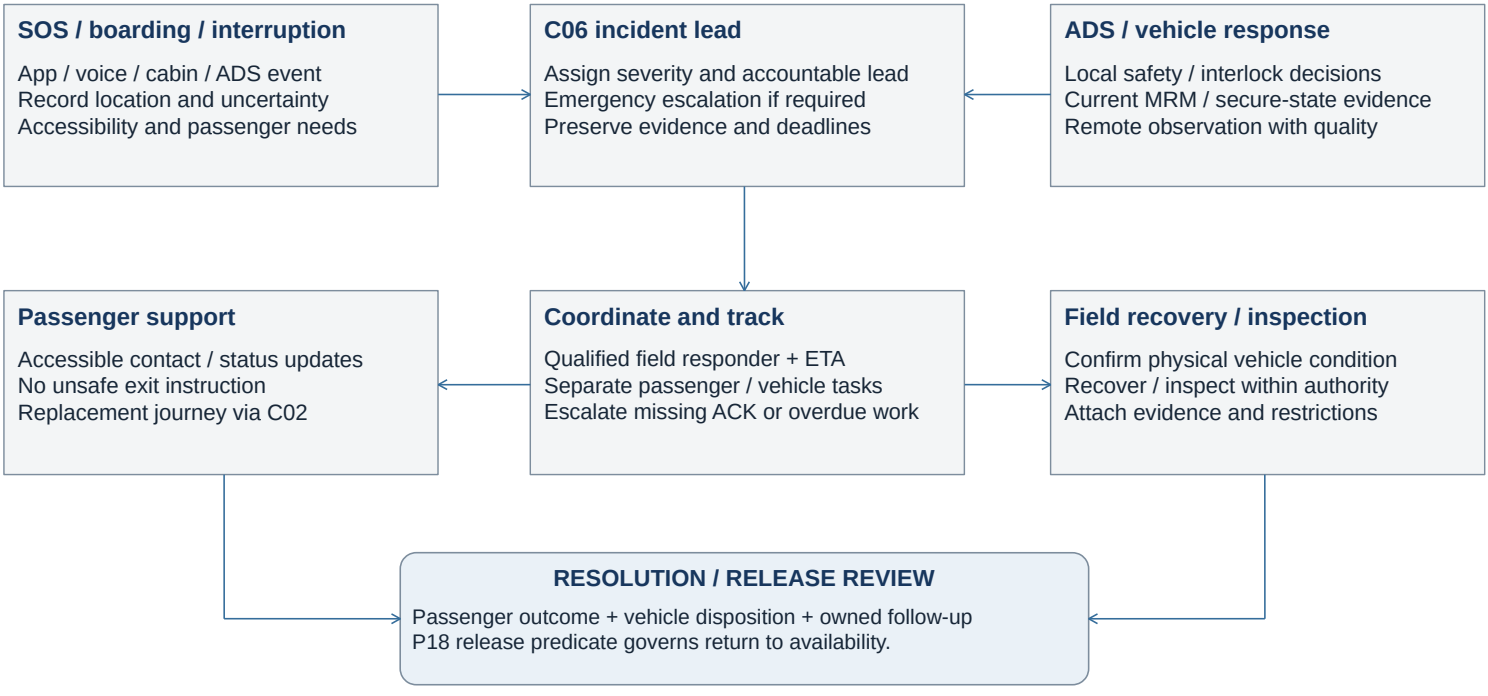
State / permitted actions

ADS mode / ODD / MRM / MRC
Motion, chassis, door, energy
Per-signal age and validity
Assess / advise / stop request
Passenger / field contact
Disabled action shows reason
Scope confirmation before submit

Action and handover record

Vehicle + mission + case revision + lease epoch + reason + advice type + expiry + expected outcome
Action-specific guards use evidence and vehicle authority. Status queries remain available with stale data.
Handover requires acceptance of pending commands, passenger status, field ETA and next deadline.
Expired / unacknowledged ownership triggers reassignment, not silent abandonment.
Relief capacity and supervisor escalation handle concurrent incidents and a failed seat.

14.2 Incident workflow with parallel vehicle and passenger outcomes



14.3 Workload and scenario acceptance

Staffing input: $A = \Sigma(\lambda_k \times E[S_k])$ in matched time units. A/m is a mean-load screen, not a service-level guarantee. Test bursts, breaks, common-cause alarms, handover and centre loss. Include field-team travel and replacement-trip capacity separately. Record detection, ACK, resolution, overdue cases and failed recovery. Restrict new commitments when response capacity is exhausted. P27 defines service-mode and return-to-service gates. Every open branch retains an owner and deadline.

Action-specific guards and time validity

RM-10 / RM-12. Permissions depend on action class. Read-only recovery and emergency escalation remain possible when telemetry is stale. Motion-enabling advice requires current contextual evidence.

15.1 Action class determines its evidence and authority requirements

Read / escalate

Authenticated status or history query
Allowed while state is UNKNOWN
No movement or allocation side effect

→

Request / advise

Current owner, fence and capability
Action-specific scene/state evidence
Stop request uses vehicle contract

→

Onboard decision

Accept/reject with explicit reason
Local ODD and safety checks prevail
No remote steering/throttle baseline

15.2 Proposed command and immutable outcome envelopes

RemoteActionRequest

tenantId, vehicleId, missionId?, commandId
caseId?, expectedCaseRev?, actionType, parameters
capabilityRev, sceneRef?, expectedVehicleStateRev?
operatorId, qualificationRef, authorityEpoch
issuedAt, expiresAt, nonce, payloadHash
reasonCode, evidenceRefs[], approvalRef?

Required fields follow a versioned action schema.
CONTEXT_ADVICE / STOP_REQUEST /
MISSION_CHANGE_REQUEST / STATUS_QUERY
Same ID with changed payload is rejected.

CommandAuditEvent

tenantId, eventId, commandId, correlationId
producerId, bootEpoch, streamId, sequence
observedAt, receivedAt, clock-bound metadata
RECEIVED / ACCEPTED / REJECTED / EXECUTING /
COMPLETED / ABORTED / EXPIRED + reason
stateRevision, evidenceRefs, policyVersion, actor

UNKNOWN belongs to the observer, not the ADS.
Keep raw event + validation result append-only.
Missing terminal evidence stays unresolved.
Expiry blocks new acceptance, not reconciliation.

15.3 Freshness uses bounded clock error and future-time rejection

$\Delta_i = t_{now} - t_{observed,i}$ $U_i = u_{source,i} + u_{observer}$
$ageUpper_i = \max(0, \Delta_i + U_i)$
$timeOK_i = boundsKnown \wedge clocksValid \wedge (\Delta_i \geq -U_i) \wedge (ageUpper_i \leq F_i)$
$valid_i = identityOK \wedge schemaOK \wedge sourceOK \wedge epoch/sequenceOK \wedge timeOK_i$
All terms use seconds. U includes synchronisation error and drift since the last trusted update.
No finite clock bound means time validity is UNKNOWN. Local timers use a monotonic clock.

Response and execution deadlines

$D_{remote} = D_{detect} + D_{queue} + D_{ack} + D_{assess} + D_{transmit} + D_{accept}$. Budget is operational. Do not use it to justify onboard safety deadlines. Expiry uses the receiver's validated time bound; clock uncertainty rejects new time-critical acceptance.

Ownership and release parameters

ADS owner: local link-loss/fallback, action capability, fence and journal persistence. Platform: F_i, clock bounds, sampling, loss, retry and retention. Operations: response class, seat/field coverage and escalation. Each value needs a unit, rationale and boundary test.

15.4 Audit scope and evidence lifecycle

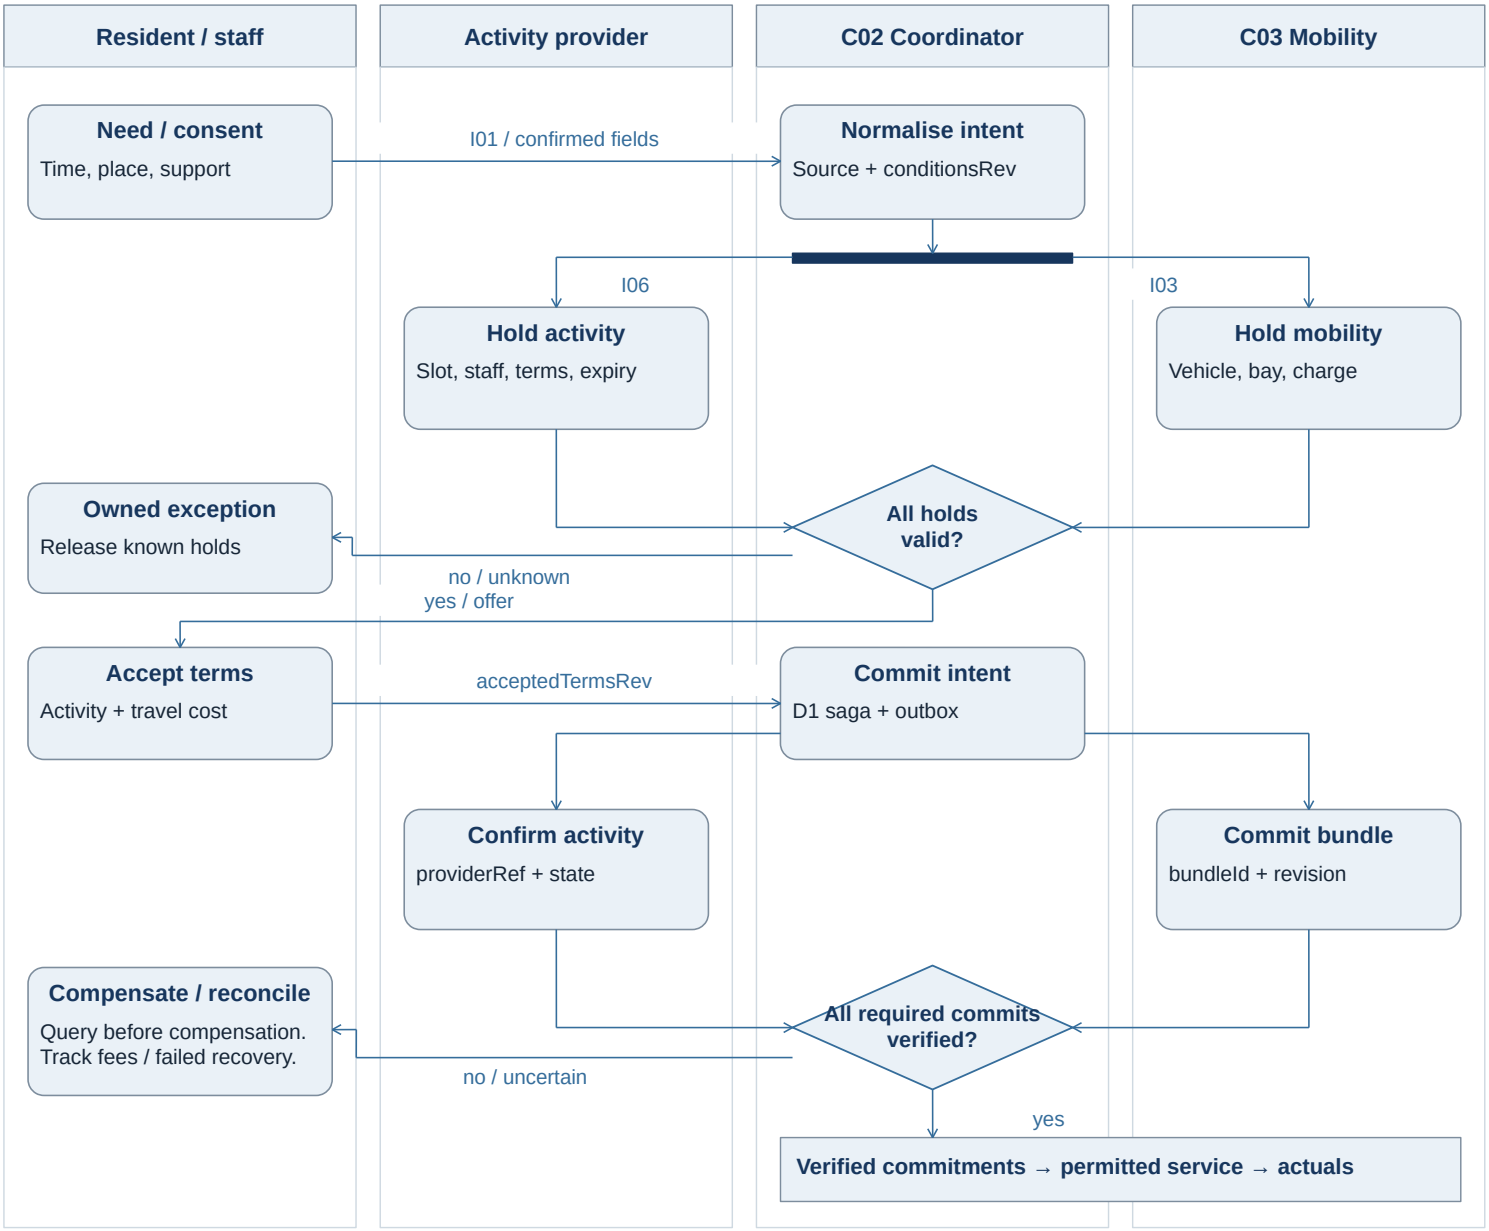
Purpose-limited capture, scoped access, encryption, retained incident evidence and read-only replay. AI may draft or summarise within task scope; it cannot independently authorise ADS movement. Test cross-tenant access, expired authority, future timestamps, unknown clock bounds and duplicate effects.

Desk-reviewed architecture proposal; implementation, performance and safety validation remain pending. Public references support terminology and patterns, not this system's certification. Sources in notes: SAE J3016 baseline, ISO 34503, NHTSA, Waymo, NPA, AWS, Microsoft, Hazelcast, PostgreSQL, RFC 5905 and calendar API documentation.

Activity saga, service workflow and AI authority

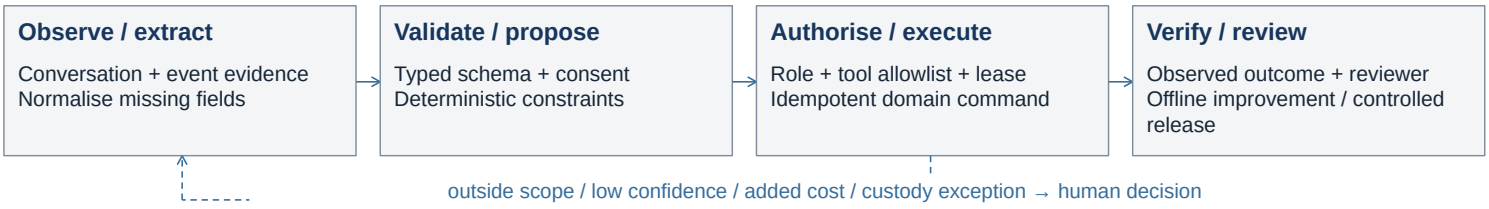
Care outings, after-school care and community events share coordination logic. Each provider retains service responsibility, staffing and handover obligations.

16.1 Swimlane activity flow with parallel holds and partial-failure compensation



Service guards: care = assistance + care-staff capacity; school care = guardian authority + verified handover; community event = venue/stop + host capacity. ETA changes re-enter the saga with a new revision. Provider and transport reservations are not one distributed transaction.

16.2 AI assistance as a constrained tool pipeline



Release stages: shadow comparison → staff-approved actions → bounded automatic tasks
Independent guards enforce consent and accepted terms. Model updates require offline evaluation and versioned release.

Calendar changes produce proposed intent, never auto-booking by default. Sync deletion/resync must not erase D1 bookings. Commercial APIs need access and hold/commit/cancel support; otherwise use an owned manual workflow.

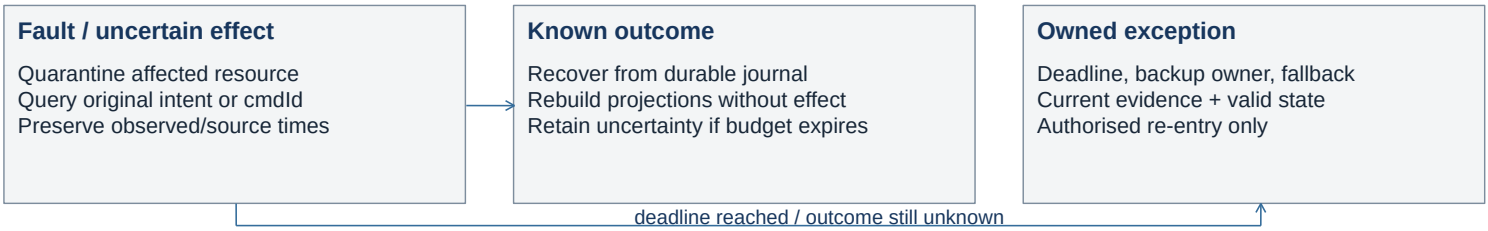
Verification coverage and release dependencies

RM-11. Remote-operation requirements join the PaaS acceptance plan. The cases below are planned tests, not passed results. Release requires reproducible evidence and operating readiness.

17.1 Fault injection and observable remote-operation invariants

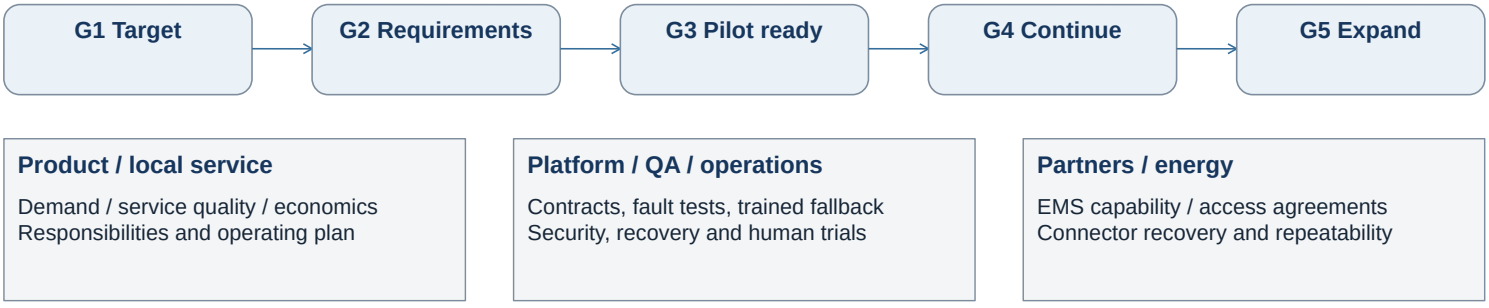
T-R01 / RM-03–04	Delay, reorder, clock jump, reboot	Stale/future data blocks motion; query stays usable.
T-R02 / RM-01,05	ODD limit / sensor degradation	Correct response, no remote extension of ODD.
T-R03 / RM-06,10	Expired / changed-scene advice	Reject with reason and renew context if required.
T-R04 / P13,RM-07	Lose ACK / replay command ID	No duplicate effect; lost result remains reconciled.
T-R05 / RM-02,07	Lose telemetry / media / voice	Independent quality flags; local policy + escalation.
T-R06 / RM-02,10	Centre partition / old operator	Fence ACK + reconcile prior accepted effects.
T-R07 / RM-05,09	SOS + MRM failure / no field team	Emergency and passenger work remain owned.
T-R08 / RM-04,08	Alarm storm / unauthorised AI	No hidden P0; deny AI motion-authorisation.

17.2 Shared PaaS invariants and fault recovery



Add roadside clock/calibration, coverage, false association, reboot and replay tests (P15–P16). P25–P28 extend acceptance to peak demand, spatial imbalance, finite hub capacity, replenishment queues and combined failures. Retain settlement regression tests.

17.3 Evidence chain, milestones and parallel workstreams



G3 evidence: scenario/ODD slice, initial state, revisions, injected fault, trace and oracle result, tested in SIL/HIL/operator/vehicle trials as applicable. G4 uses observed outcomes. G5 adds repeatable onboarding. Record owner, residual risk/restriction and signed decision at each gate.

Service capacity and geographic coverage

Operating feasibility depends on time and location. Fleet and hub totals become usable capacity only when vehicle readiness, travel time and the full service cycle align.

18.1 Demand, reachable supply and commitment

Demand by zone and time

OD requests / time window
Booking lead time and peak shape
Party size and accessibility needs
Customer cancellations recorded

→

Feasible supply

Current position and ready time
Vehicle / stop / ODD compatibility
Pickup and arrival time windows
Energy and next-task constraints

→

Service commitment

Hold the complete resource bundle
Confirm only after verified commit
Return a binding reason if infeasible
Retain unmet demand in reporting

↓

C03 evaluation at each planning interval

$t_{ready,v} + travelTime(v, origin, departure\ time) \leq latestPickup$
Next tasks include passenger travel, repositioning, turnaround and replenishment.
A reservation lead time enables physical pre-positioning only when time and energy remain feasible.

18.2 Fleet conservation and workload screen

Ready	Committed	In service	Replenishing	Unavailable
Assignable now	Held / pre-position	Pickup / passenger	Queue / energy	Fault / inspection

$$N_{total} = N_{ready} + N_{committed} + N_{service} + N_{replenishing} + N_{unavailable}$$

Necessary capacity screen

λ [requests/h] × E[T_{cycle}] [h/request] ≤ $N_{service\text{-}capable}$
T_{cycle} includes occupied and empty travel, turnaround and attributable replenishment time.
Evaluate with mutually exclusive states. Do not deduct the same downtime twice.

Spatial and peak-time check

Average capacity does not prove pickup feasibility.
Test demand bursts and directional imbalance by zone.
Hold operational reserve for faults and recovery.
Intercity transfers require explicit travel and authority.

18.3 Hub function, allocation and service outcome

Hub resources

Parking bays and pickup space
Connectors / dispensers / site power
Cleaning and inspection positions
Qualified staff and opening hours

→

Finite allocation

Use time intervals per resource
Sum pooled demand against limits
Reserve only accessible capacity
Propagate queues and outage derating

→

Observable outcome

Dispatch / pickup / completion
Unserved and delayed requests
Queue and recovery workload
Energy and cost per completed trip

Planning record: zone / time interval, OD input, initial vehicle positions, hub resources, availability policy and service windows. Compare feasible alternatives under the same demand trace. A hub count alone does not specify throughput or coverage.

Replenishment and departure readiness

The operating plan reserves replenishment capacity within a vehicle's duty cycle. Daily energy sufficiency and on-time departure are separate acceptance conditions.

19.1 Vehicle cycle with finite replenishment resources

Arrive
Observed energy
Arrival time

→

Queue
Bay / connector
Staff eligibility

→

Replenish
Power or fuel
Actual delivery

→

Turnaround
Clean / inspect
Required evidence

→

Release
Energy + tasks
Next departure

Departure constraint

$t_{arrival} + t_{queue} + t_{replenish} + t_{required_work} \le t_{departure}$
Only explicitly approved concurrent tasks may overlap. Travel to a different site consumes time and energy.
A late or uncertain delivery triggers re-planning before the customer promise becomes infeasible.

19.2 Electrical and fuel branches use distinct units

Electrical charging

$E_{depart} = E_{arrive} + \eta_{ch} \int P_{AC}(t)\,dt - E_{aux}$
$E_{next_trip} + E_{reserve} \le E_{depart} \le E_{usable,max}$
P_{AC} in kW, time in hours, energy in kWh.
Use charger / battery limits and charging efficiency.
Energy plan ACK does not prove delivered energy.

Fuel / hydrogen replenishment

$m_{depart} = m_{arrive} + m_{delivered} - m_{consumed}$
$m_{next_trip} + m_{reserve} \le m_{depart} \le m_{usable,max}$
Mass in kg, flow in kg/min, duration in minutes.
Bound usable stock, dispenser flow and turnaround.
Fuel-to-energy conversion needs an explicit model.

SITE LIMITS AT EVERY TIME INTERVAL

$\Sigma P_{charge,i}(t) \le P_{site,available}(t)$ and occupied connectors ≤ serviceable connectors
$P_{site,available}$ reflects building load, generation, storage and connection limits at the same boundary.
Fuel stock remains non-negative after deliveries and withdrawals. One hub can contain several resources.

19.3 Shortage propagation and recovery

Detect
Power derating / site outage
Low stock / supply delay
Queue growth / delayed arrival
Stale capability or meter data

→

Re-plan
Recalculate departure readiness
Use feasible alternative site / vehicle
Hold affected assignments
Revise ETA and notify service owner

→

Reconcile and release
Observe actual delivered quantity
Complete inspection and required work
Confirm current resource revisions
Release only verified capacity

Acceptance evidence: arrival / departure energy, measured delivery, queue duration, site-limit trace and missed departures. Test concurrent arrivals, one resource unavailable, reduced power and delayed fuel delivery. Report reserve violations and unmet trips even when aggregate daily energy balances.

Degraded operation and recovery capacity

Operational continuity requires bounded service reductions, passenger support and a verified return to service. Recovery uses the same finite fleet, site and staffing resources.

20.1 Failure propagation across the service system

Demand pressure

Peak arrivals / directional imbalance
Longer trips / slow turnaround

Supply loss

Vehicle fault / site outage / energy shortfall
Communication loss / staff absence

↓

Impact assessment

Identify affected bookings and intervals. Recompute reachable vehicles, usable hubs and qualified response capacity.
Preserve already accepted commitments and unknown physical effects. Record unmet demand and passenger exposure.

↓

NORMAL

Service promise feasible
Capacity and evidence current

→

CONSTRAINED

Restrict new offers by zone / time
Protect accepted journeys

→

SUSPENDED

Stop affected new commitments
Maintain support and local fallback

Transitions use configured triggers, persistence and owner approval. Recovery thresholds include hysteresis to avoid repeated entry / exit. Local ADS safety decisions remain independent of this service-mode model.

20.2 Parallel passenger, asset and capacity recovery

Passenger outcome

Contact through accessible channel
Explain revised ETA / alternatives
Arrange feasible replacement journey
Track handover and resolution
Own compensation / refund decision

Vehicle / site outcome

Retain fault and command history
Assign qualified field responder
Observe secure state and clearance
Inspect and record restrictions
Reconcile unresolved physical effects

Resource outcome

Reserve replacement and field capacity
Include responder travel and queue
Recalculate affected bookings
Track backlog age and deadlines
Escalate resource exhaustion

20.3 Return-to-service gate and staffing condition

Return-to-service gate

Cause addressed + fresh evidence + authorised release + feasible next mission + passenger outcome resolved.
Connectivity restoration alone cannot release capacity.
Observe sustained recovery before widening offers.

Response capacity

$A = \Sigma \lambda_k E[S_k]$ in matched time units.
$A <$ qualified staffed seats is only a mean-load screen.
Test bursts, common-cause alarms, breaks and travel.
Every open case needs an owner, deadline and backup.
Keep field response separate from remote assistance.

Record trigger, affected scope, admission restriction, accountable owner, next review and recovery evidence for each degraded period.

Operating feasibility and verification evidence

A configuration qualifies only within its tested demand, geography and operating conditions. Service quality, resource limits, recovery and economics require separate evidence.

21.1 Measurement chain and explicit denominators

Requests
Unique service intents

→

Dispatch
Verified assignment

→

Pickup
Physical arrival

→

Completion
Verified trip outcome

Service quality

Dispatch success = assigned / eligible dispatch requests.
Keep capacity and admission rejects in the denominator.
Demand fulfillment = completed / eligible requests.
Report customer cancellations separately by reason.

Pickup and utilization

Wait starts at request time for immediate demand, or agreed pickup time for a reservation. Report p50 / p95 with the observed-arrival count and missing outcomes.
Show on-time pickups / pickups due, including failures.
Separate available hours from scheduled fleet hours.

Count each request once at the reporting cut-off: completed, cancelled, rejected, expired, unresolved or still active. Preserve event history separately. Request cohorts and deadlines determine metric eligibility. Show exclusions, censoring and missing evidence.

21.2 Reproducible scenario design and pass conditions

Freeze inputs

OD / hourly arrivals and trip lengths
Booking share and lead-time profile
Initial fleet positions / availability
Hub, energy and response limits
Geography and parameter revisions

→

Exercise conditions

Normal demand and forecast error
Peak and directional concentration
Vehicle / hub / energy capacity loss
Link outage and support overload
Combined faults and recovery ramp

→

Evaluate each case

Service targets by zone and time
Resource and energy conservation
Queue age and passenger recovery
Missing / unresolved outcomes
Binding constraint and owner

PASS RECORD

State target, denominator, time window and acceptance rule before running. Use repeatable demand traces or recorded seeds, then independent trials. Preserve every run and failure. Finite synthetic tests establish a tested envelope only. Pilot evidence determines field readiness.

21.3 Capacity decision and business consequence

Feasible alternatives

Compare fleet, hub placement and service policy only after all required operating conditions pass.
Rank cost within the feasible tested set. If none qualifies, report the binding constraint and re-plan.

Economic reconciliation

Use completed trips and measured resource use.
Include empty travel, energy, support and recovery.
Separate CAPEX, annual operating cost and cash flow.
Avoid double-counting integrated hub infrastructure.

Release evidence: input revision, scenario / seed, execution trace, metric definitions, pass / fail by condition, residual restriction and accountable decision. Operational feasibility does not by itself establish project profitability or safety certification.